

Supervised Learning of Behaviors

CS 285

Instructor: Sergey Levine
UC Berkeley



Terminology & notation



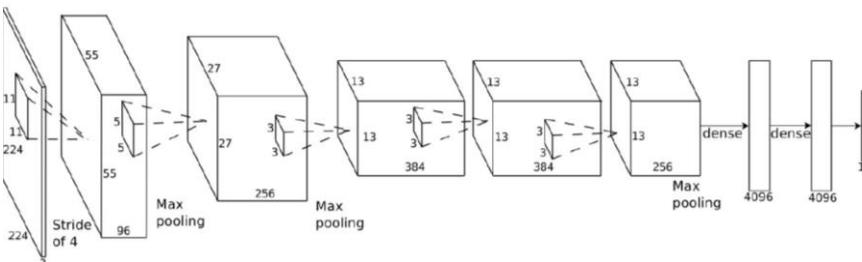
\mathbf{o}_t



\mathbf{s}_t – state

\mathbf{o}_t – observation

\mathbf{a}_t – action



$$\pi_{\theta}(\mathbf{a} | \mathbf{o})$$



\mathbf{a}_t



$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)$ – policy

$\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)$ – policy (fully observed)



\mathbf{o}_t – observation

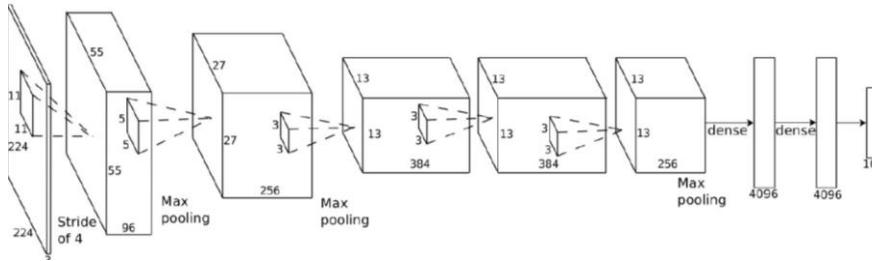


\mathbf{s}_t – state

Terminology & notation



\mathbf{o}_t



$\pi_\theta(\mathbf{a}_t | \mathbf{o}_t)$



\mathbf{a}_t

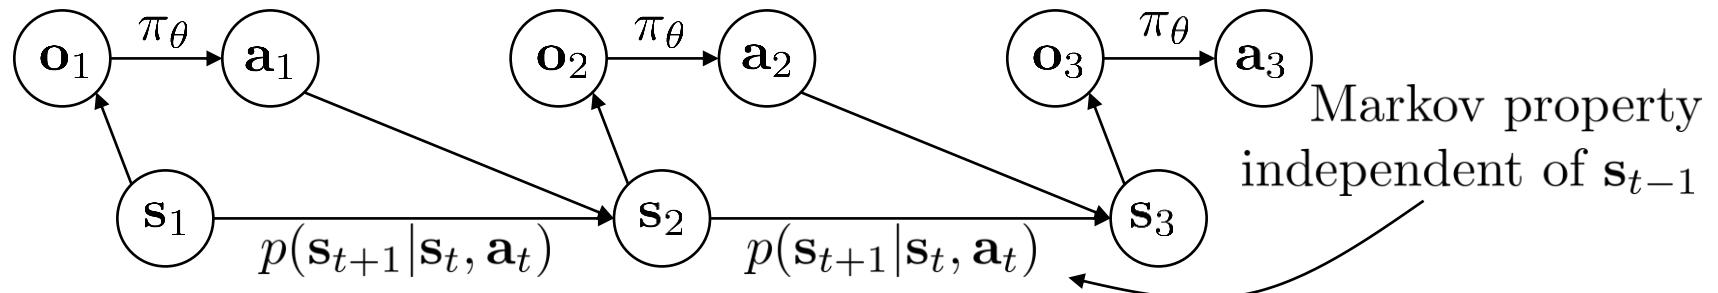
\mathbf{s}_t – state

\mathbf{o}_t – observation

\mathbf{a}_t – action

$\pi_\theta(\mathbf{a}_t | \mathbf{o}_t)$ – policy

$\pi_\theta(\mathbf{a}_t | \mathbf{s}_t)$ – policy (fully observed)



Aside: notation

s_t – state

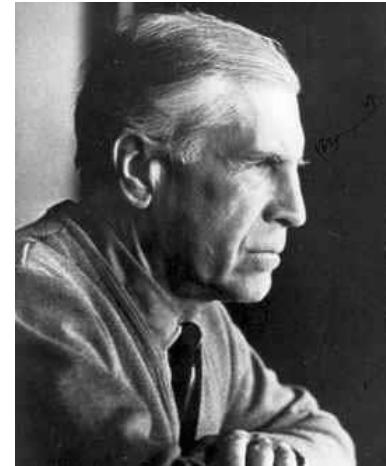
a_t – action

x_t – state

u_t – action управление

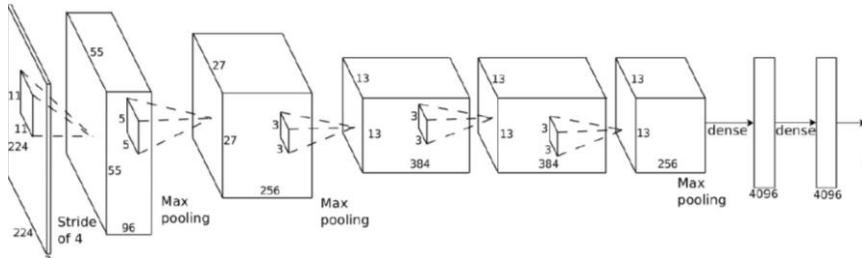


Richard Bellman

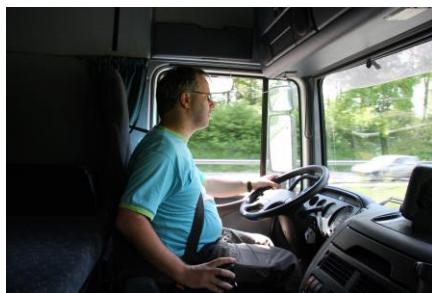


Lev Pontryagin

Imitation Learning



$$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)$$



\mathbf{o}_t



training
data



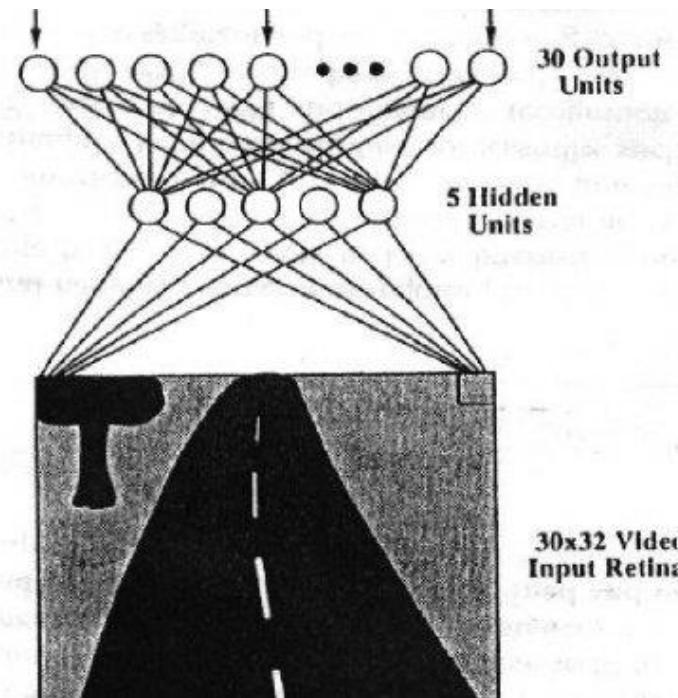
supervised
learning

$$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)$$

behavioral cloning

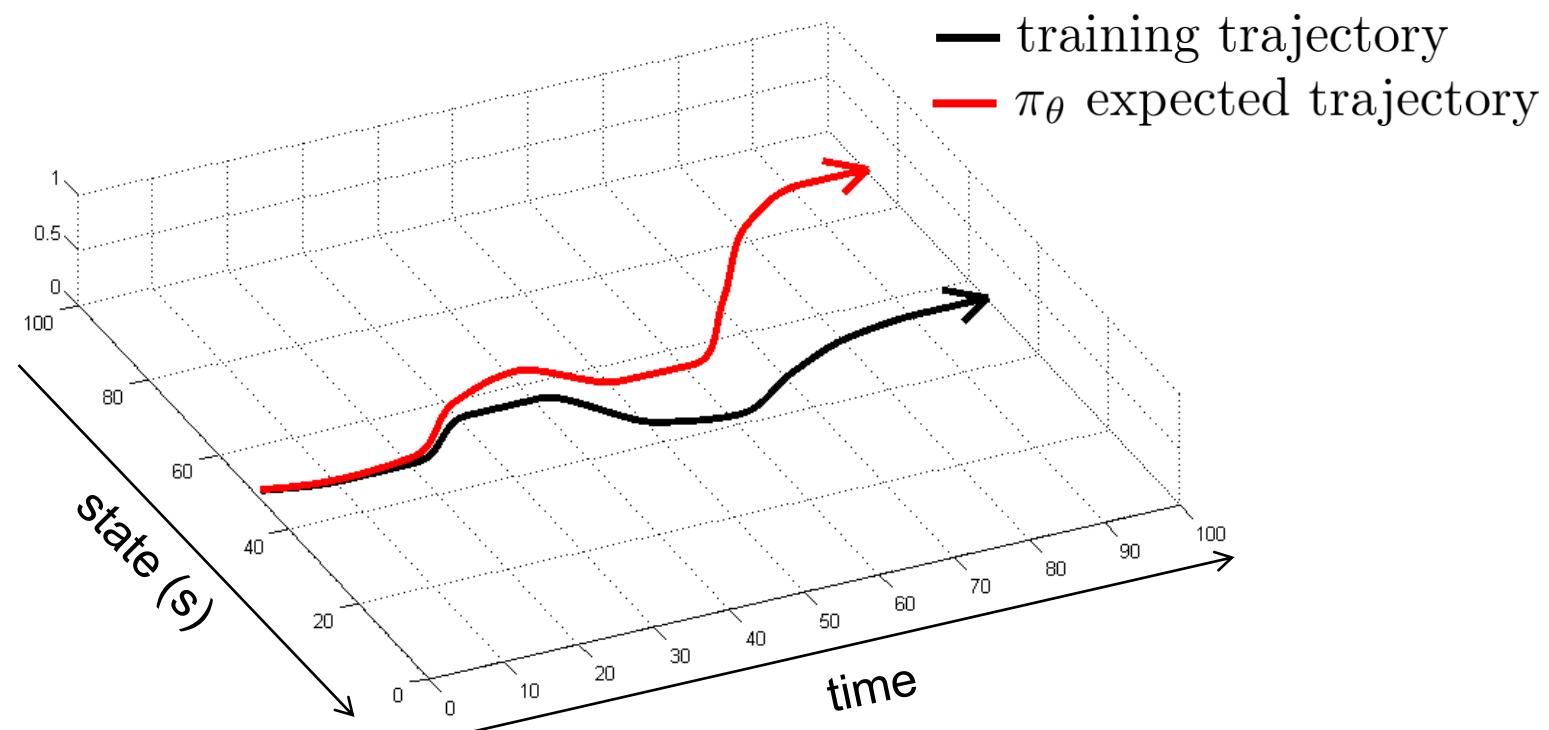
The original deep imitation learning system

ALVINN: Autonomous Land Vehicle In a Neural Network
1989



Does it work?

No!

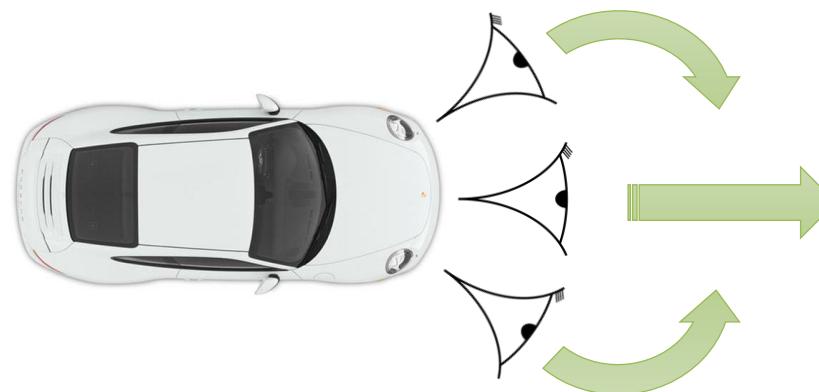
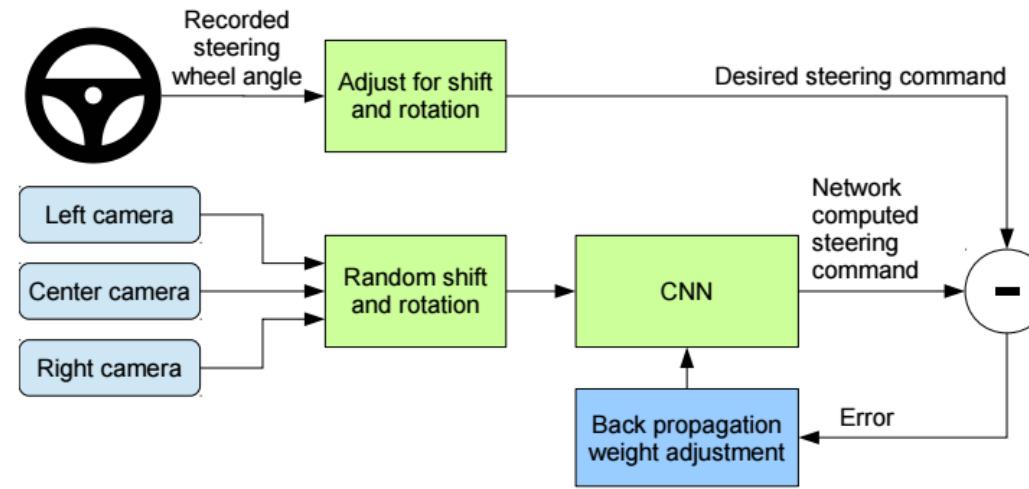


Does it work?

Yes!

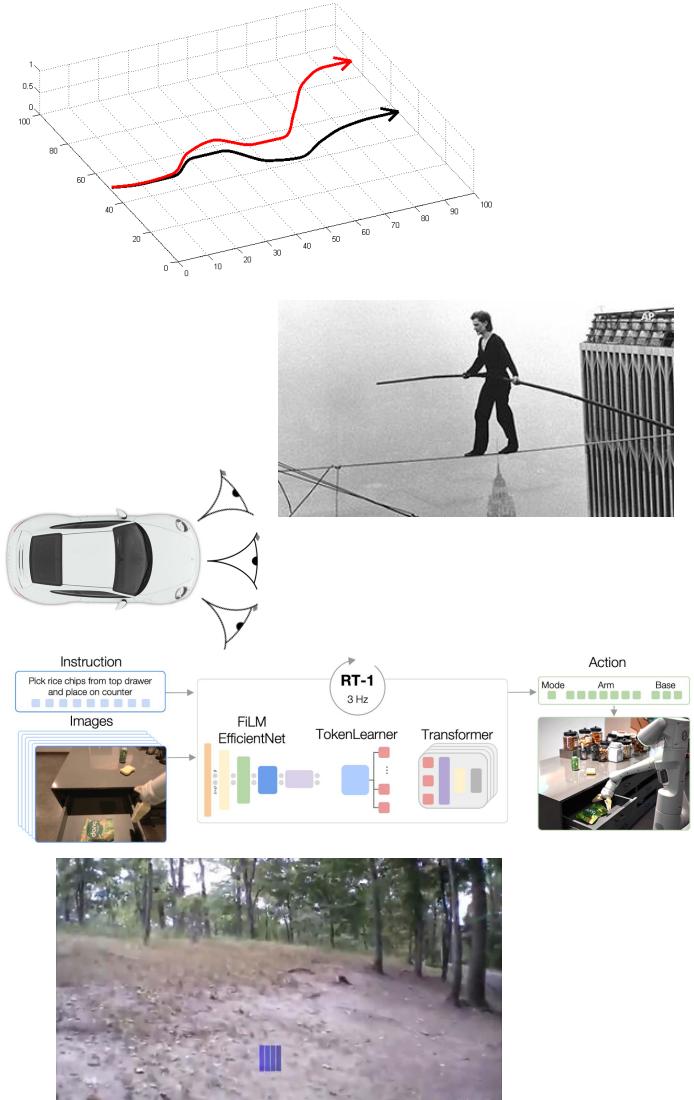
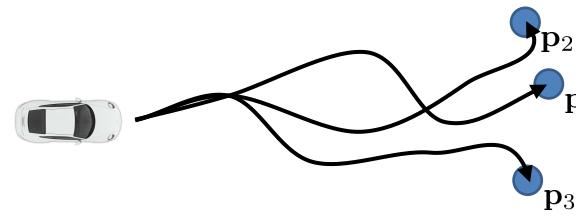


Why did that work?



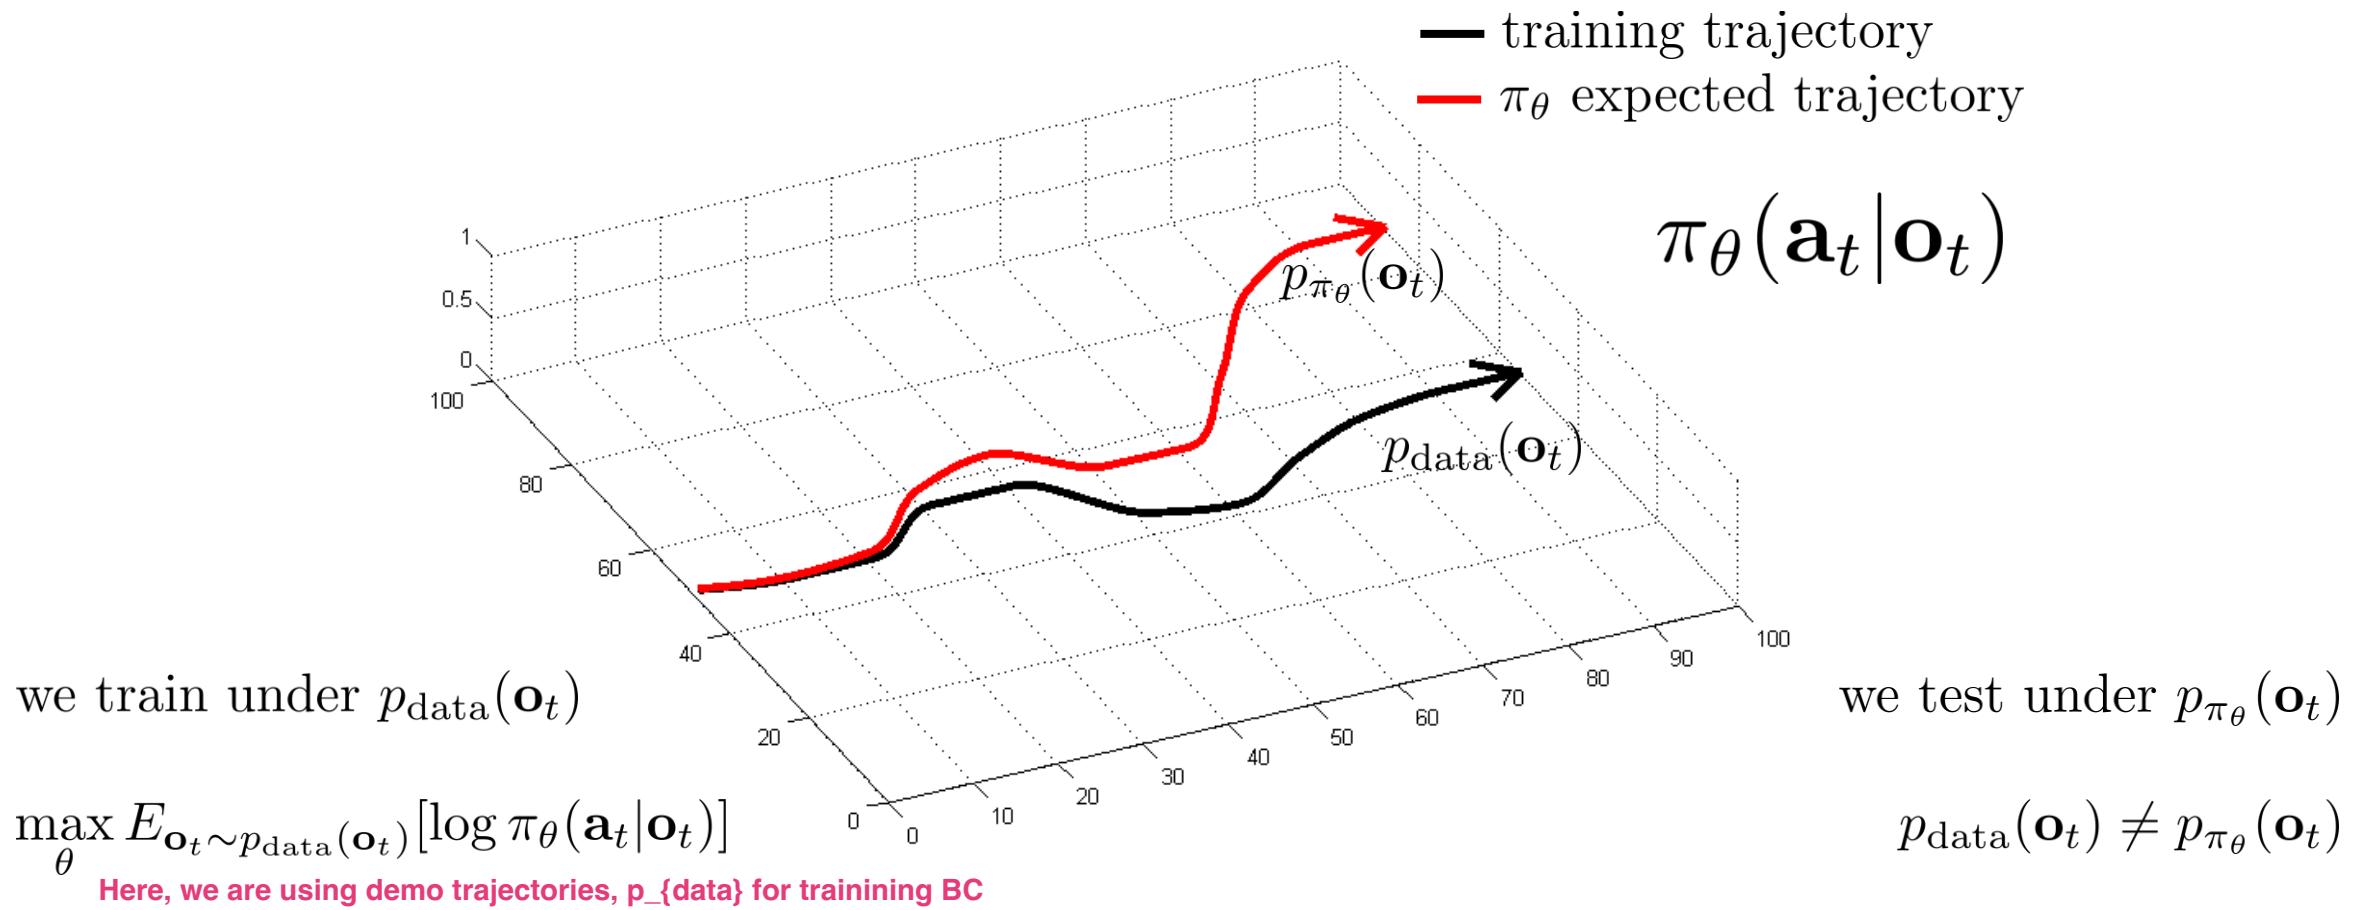
The moral of the story, and a list of ideas

- Imitation learning via behavioral cloning is not guaranteed to work
 - This is **different** from supervised learning
 - The reason: i.i.d. assumption does not hold!
- We can formalize **why** this is and do a bit of theory
- We can address the problem in a few ways:
 - Be smart about how we collect (and augment) our data
 - Use very powerful models that make very few mistakes
 - Use multi-task learning
 - Change the algorithm (DAgger)

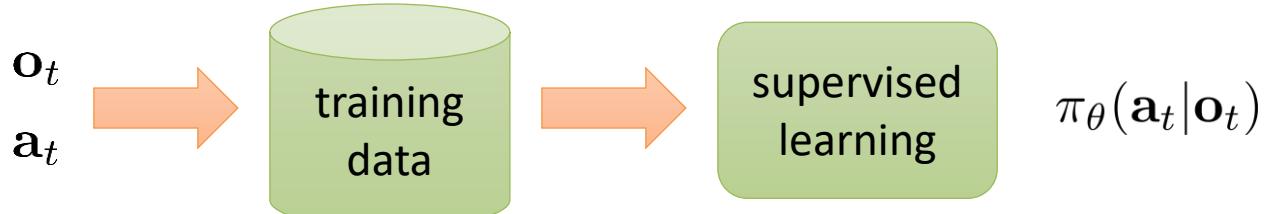


Why does behavioral cloning fail?
A bit of theory

The distributional shift problem



Let's define more precisely what we want



What makes a learned $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ good or bad?

LHS & RHS both are stochastic term, LHS follows learned policy, RHS the demo policy

$$c(\mathbf{s}_t, \mathbf{a}_t) = \begin{cases} 0 & \text{if } \mathbf{a}_t = \pi^*(\mathbf{s}_t) \\ 1 & \text{otherwise} \end{cases}$$

Goal: minimize $E_{\mathbf{s}_t \sim p_{\pi_\theta}(\mathbf{s}_t)}[c(\mathbf{s}_t, \mathbf{a}_t)]$

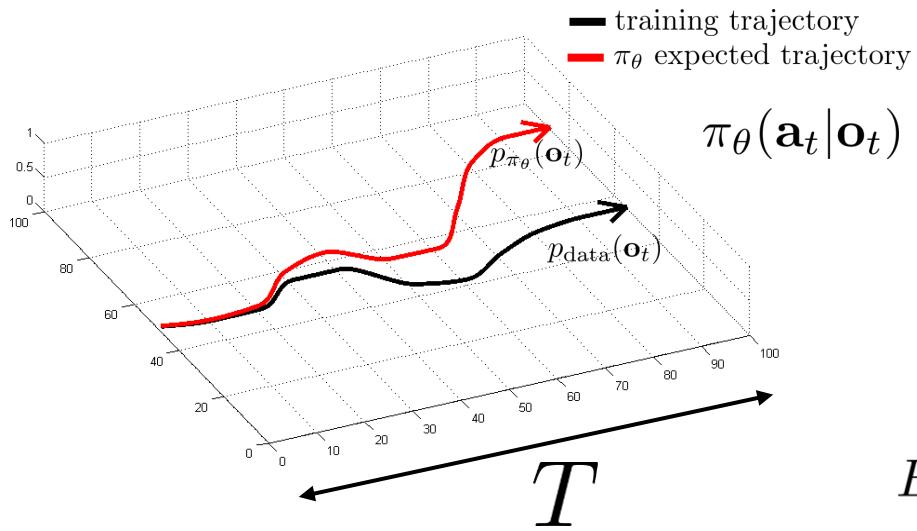
Here, we use trajectory of the learned policy

$$\cancel{\max_{\theta} E_{\mathbf{o}_t \sim p_{\text{data}}(\mathbf{o}_t)}[\log \pi_\theta(\mathbf{a}_t|\mathbf{o}_t)]}$$

Note: I started mixing up \mathbf{s} and \mathbf{o}
I warned you about that...

“Minimize the number of mistakes
the policy makes when we run it”

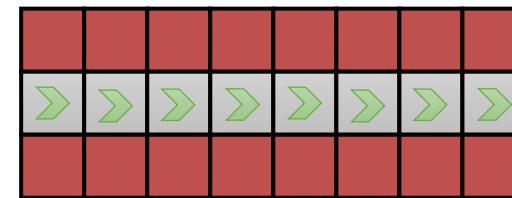
Some analysis



$$c(\mathbf{s}, \mathbf{a}) = \begin{cases} 0 & \text{if } \mathbf{a} = \pi^*(\mathbf{s}) \\ 1 & \text{otherwise} \end{cases}$$

assume: $\pi_\theta(\mathbf{a} \neq \pi^*(\mathbf{s}) | \mathbf{s}) \leq \epsilon$

for all $\mathbf{s} \in \mathcal{D}_{\text{train}}$ **This implies $E[c(\mathbf{s}, \mathbf{a})] \leq \epsilon$**



$$E \left[\sum_t c(\mathbf{s}_t, \mathbf{a}_t) \right] \leq \epsilon T + O(\epsilon T^2)$$

T terms, each $O(\epsilon T)$

More general analysis

$$c(s, a) = \begin{cases} 0 & \text{if } a = \pi^*(s) \\ 1 & \text{otherwise} \end{cases}$$

assume: $\pi_\theta(\mathbf{a} \neq \pi^\star(\mathbf{s}) | \mathbf{s}) \leq \epsilon$

for all $\mathbf{s} \in \mathcal{D}_{\text{train}}$ for $\mathbf{s} \sim p_{\text{train}}(\mathbf{s})$

actually enough for $E_{p_{\text{train}}(\mathbf{s})}[\pi_\theta(\mathbf{a} \neq \pi^*(\mathbf{s}) | \mathbf{s})] \leq \epsilon$

if $p_{\text{train}}(\mathbf{s}) \neq p_\theta(\mathbf{s})$:

$$p_{\theta}(\mathbf{s}_t) = \underbrace{(1 - \epsilon)^t p_{\text{train}}(\mathbf{s}_t)}_{\text{Correct}} + \underbrace{(1 - (1 - \epsilon)^t)) p_{\text{mistake}}(\mathbf{s}_t)}_{\text{Incorrect}}$$

probability we made no mistakes

some *other* distribution

This almost feels like we have concluded that:

Making mistake at every time-step is independent,
and the i.i.d probability of making mistake is epsilon
how strange ??

More general analysis

assume: $\pi_\theta(\mathbf{a} \neq \pi^*(\mathbf{s})|\mathbf{s}) \leq \epsilon$ Probability of making mistake on any state is bounded by epsilon | demo policy is deterministic in ref2

for all $\mathbf{s} \in \mathcal{D}_{\text{train}}$ for $\mathbf{s} \sim p_{\text{train}}(\mathbf{s})$

How is an inequality converts to a deterministic decomposition of policy probability density function ??

$$p_\theta(\mathbf{s}_t) = \underbrace{(1 - \epsilon)^t}_{\text{probability we made no mistakes}} p_{\text{train}}(\mathbf{s}_t) + \underbrace{(1 - (1 - \epsilon)^t))}_{\text{some other distribution}} p_{\text{mistake}}(\mathbf{s}_t)$$

how to show validity of p_mistake??

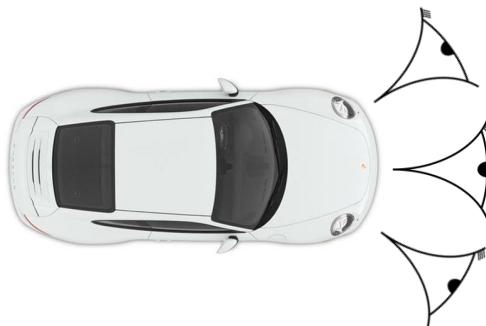
We actually don't care if this p_mistake is a valid distribution or not –
it takes value in [-1, 1]

$$|p_\theta(\mathbf{s}_t) - p_{\text{train}}(\mathbf{s}_t)| = (1 - (1 - \epsilon)^t) |p_{\text{mistake}}(\mathbf{s}_t) - p_{\text{train}}(\mathbf{s}_t)| \leq 2(1 - (1 - \epsilon)^t)$$

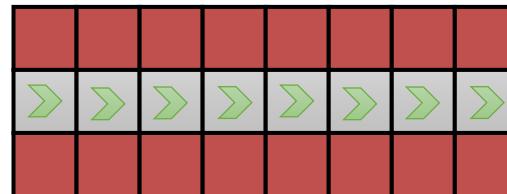
$$\text{useful identity: } (1 - \epsilon)^t \geq 1 - \epsilon t \text{ for } \epsilon \in [0, 1] \quad \leq 2\epsilon t$$

$$\begin{aligned} \sum_t E_{p_\theta(\mathbf{s}_t)}[c_t] &= \sum_t \sum_{\mathbf{s}_t} p_\theta(\mathbf{s}_t) c_t(\mathbf{s}_t) \leq \sum_t \sum_{\mathbf{s}_t} p_{\text{train}}(\mathbf{s}_t) c_t(\mathbf{s}_t) + |p_\theta(\mathbf{s}_t) - p_{\text{train}}(\mathbf{s}_t)| c_{\max} \\ &\leq \sum_t \epsilon + 2\epsilon t \\ &= O(\epsilon T^2) \end{aligned}$$

Why is this rather pessimistic?

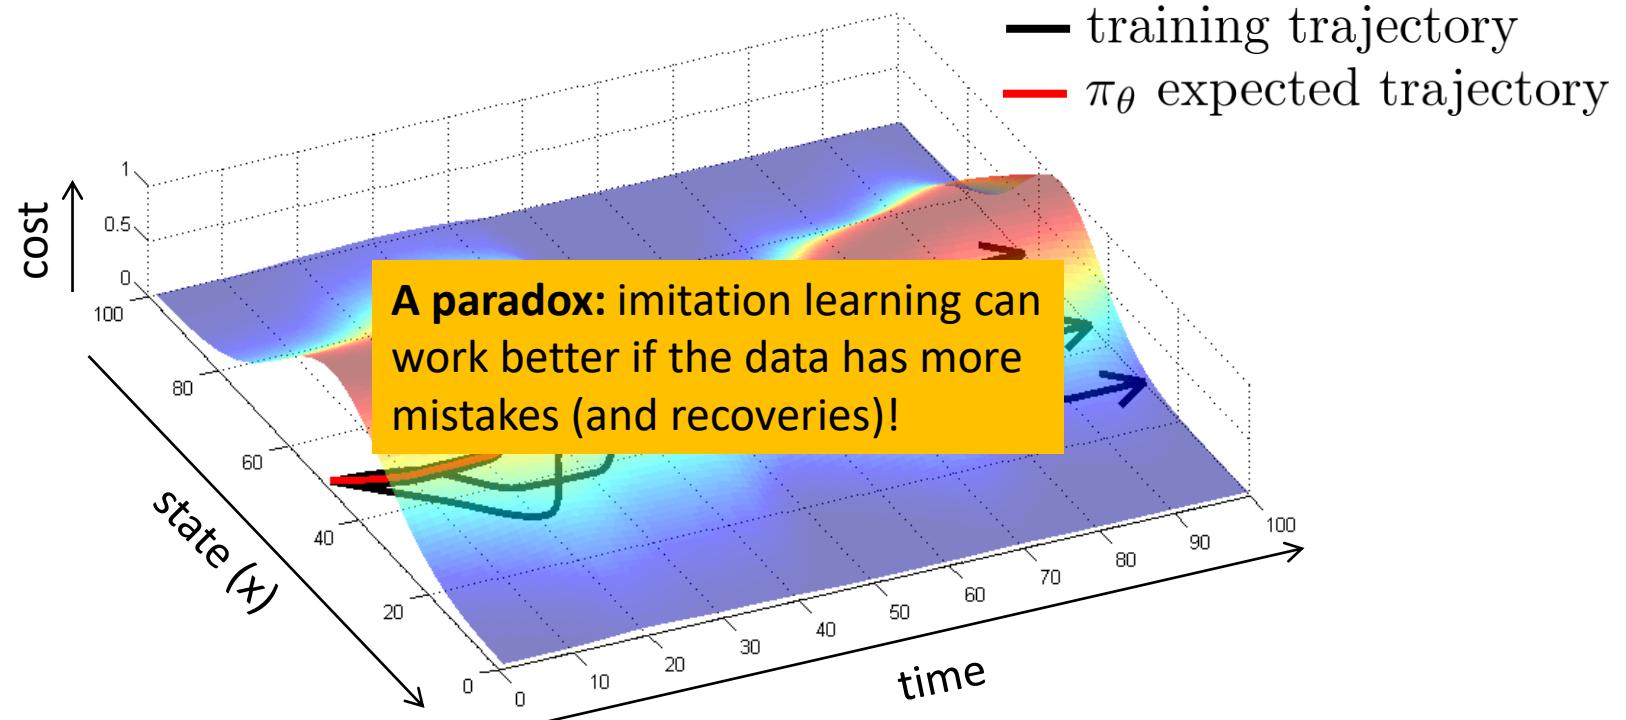


Why does this work?



In reality, we can often **recover** from mistakes

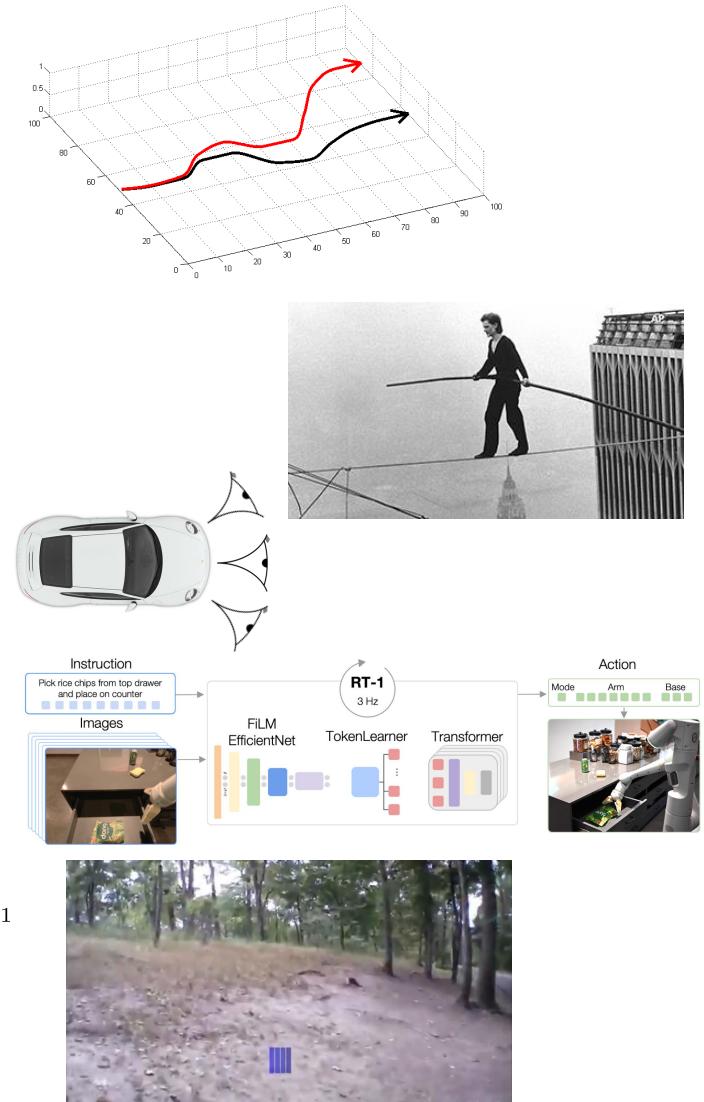
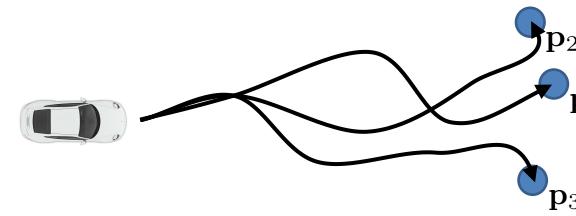
But that doesn't mean that **imitation learning** will allow us to learn how to do that!



Addressing the problem in practice

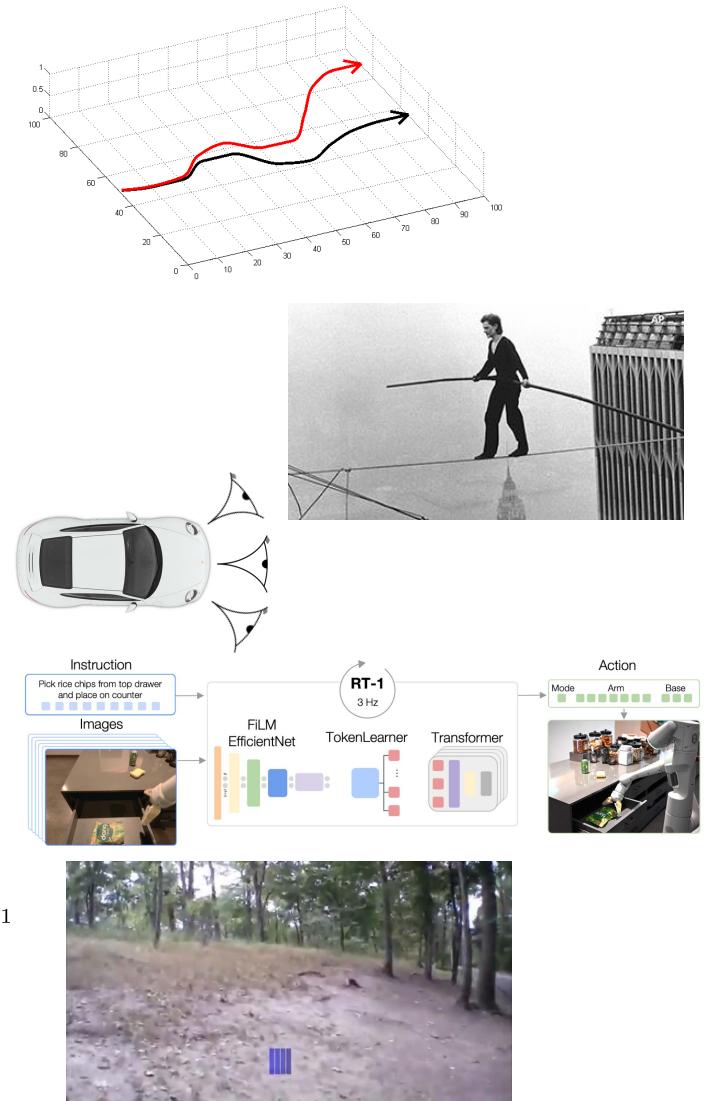
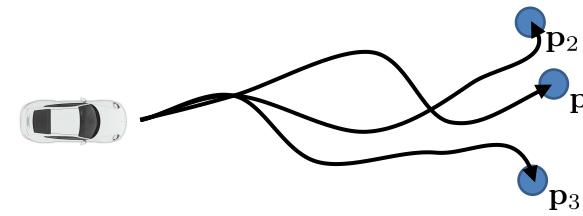
Where are we...

- Imitation learning via behavioral cloning is not guaranteed to work
 - This is **different** from supervised learning
 - The reason: i.i.d. assumption does not hold!
- We can formalize **why** this is and do a bit of theory
- We can address the problem in a few ways:
 - Be smart about how we collect (and augment) our data
 - Use very powerful models that make very few mistakes
 - Use multi-task learning
 - Change the algorithm (DAgger)

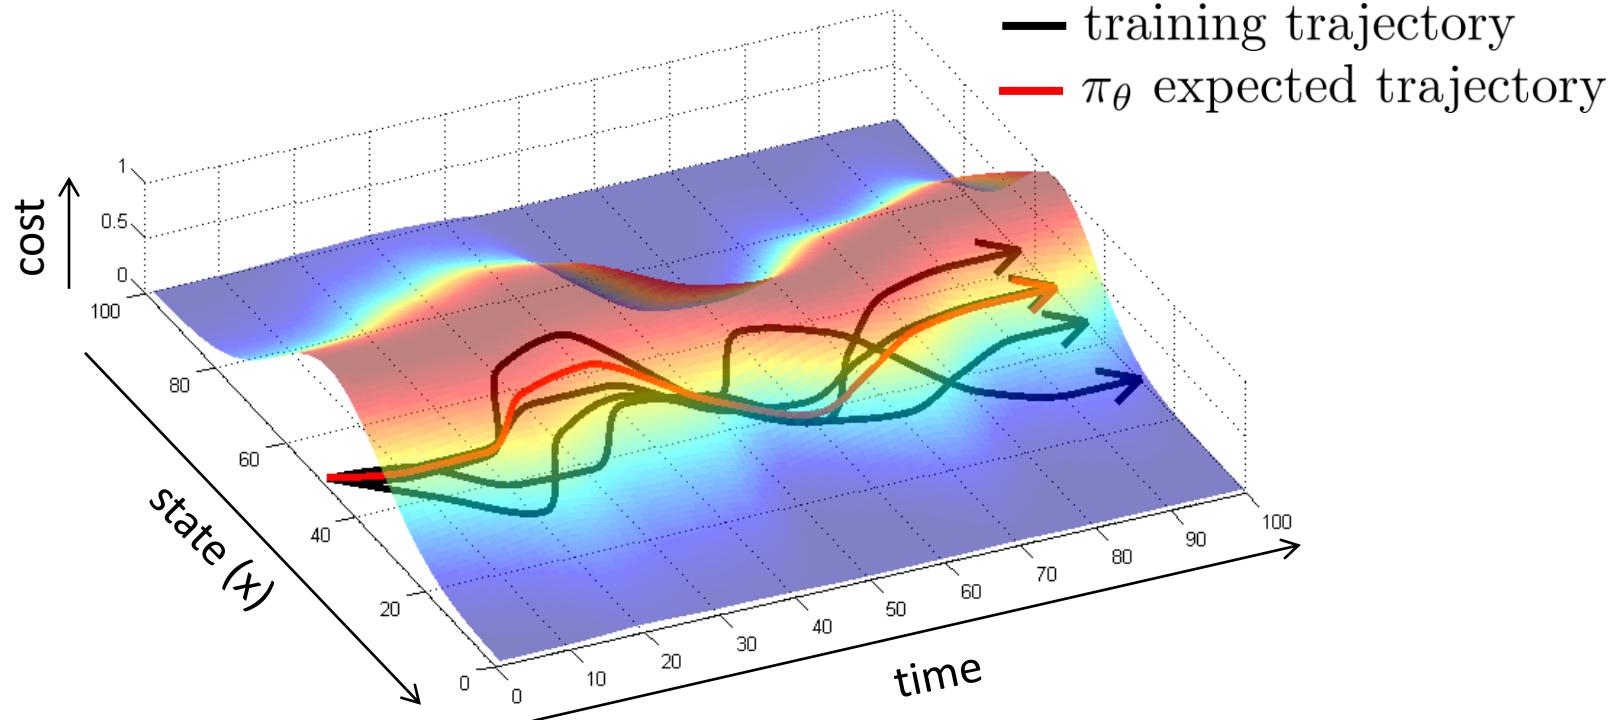


Where are we...

- Imitation learning via behavioral cloning is not guaranteed to work
 - This is **different** from supervised learning
 - The reason: i.i.d. assumption does not hold!
- We can formalize **why** this is and do a bit of theory
- We can address the problem in a few ways:
 - Be smart about how we collect (and augment) our data
 - Use very powerful models that make very few mistakes
 - Use multi-task learning
 - Change the algorithm (DAgger)



What makes behavioral cloning easy and what makes it hard?



- Intentionally add **mistakes** and **corrections**
 - The mistakes hurt, but the corrections help, often more than the mistakes hurt
- Use **data augmentation**
 - Add some “fake” data that illustrates corrections (e.g., side-facing cameras)

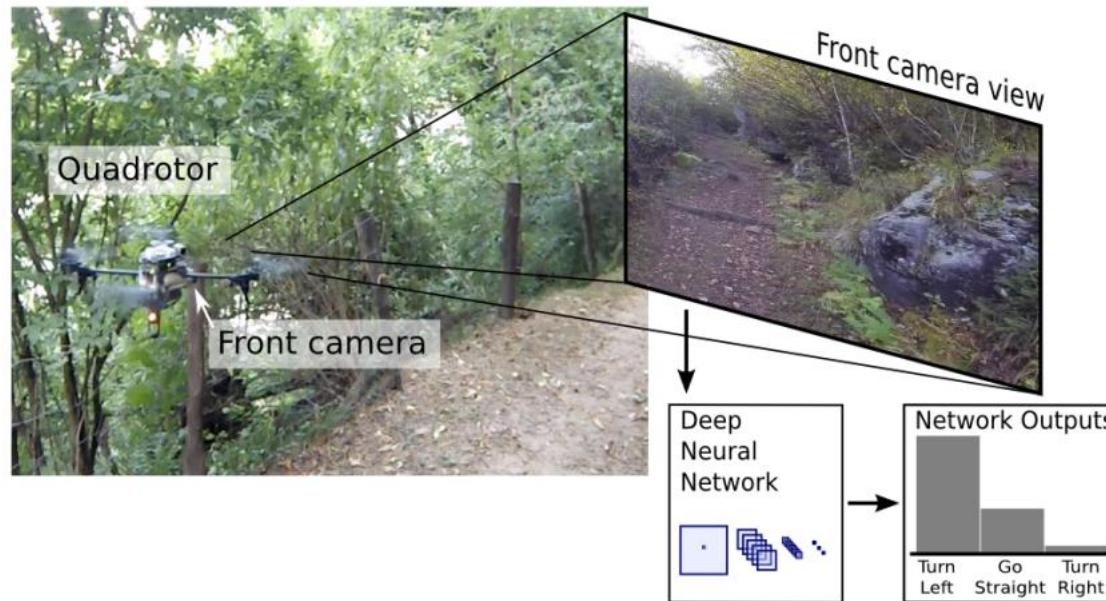
Case study 1: trail following as classification

A Machine Learning Approach to Visual Perception of Forest Trails for Mobile Robots

Alessandro Giusti¹, Jérôme Guzzi¹, Dan C. Cireşan¹, Fang-Lin He¹, Juan P. Rodríguez¹

Flavio Fontana², Matthias Faessler², Christian Forster²

Jürgen Schmidhuber¹, Gianni Di Caro¹, Davide Scaramuzza², Luca M. Gambardella¹



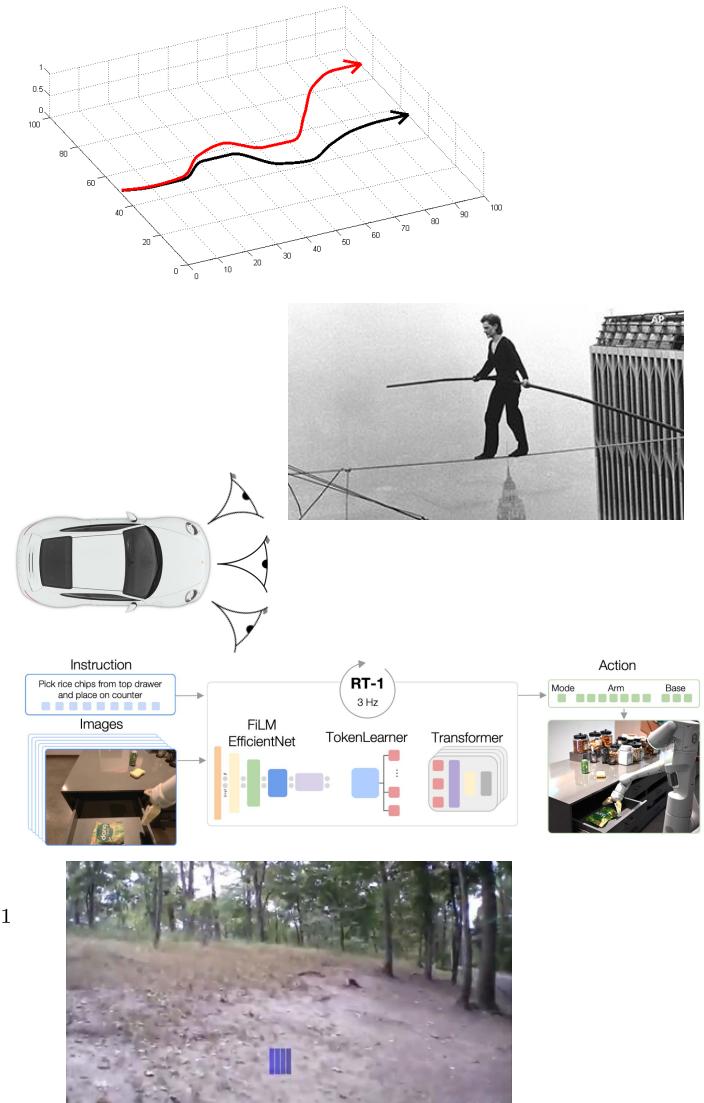
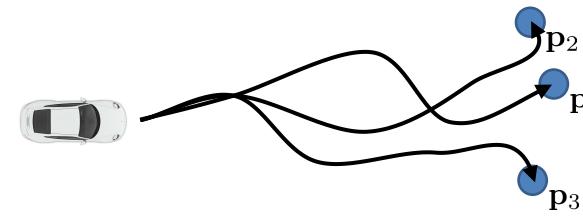
Case study 2: imitation with a cheap robot

Vision-Based Multi-Task Manipulation
for Inexpensive Robots
Using End-To-End Learning from Demonstration

Rouhollah Rahmatizadeh, Pooya Abolghasemi, Ladislau Boloni, Sergey Levine

Where are we...

- Imitation learning via behavioral cloning is not guaranteed to work
 - This is **different** from supervised learning
 - The reason: i.i.d. assumption does not hold!
- We can formalize **why** this is and do a bit of theory
- **We can address the problem in a few ways:**
 - Be smart about how we collect (and augment) our data
 - **Use very powerful models that make very few mistakes**
 - Use multi-task learning
 - Change the algorithm (Dagger)



Why might we fail to fit the expert?



1. Non-Markovian behavior
2. Multimodal behavior

$$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)$$

behavior depends only
on current observation

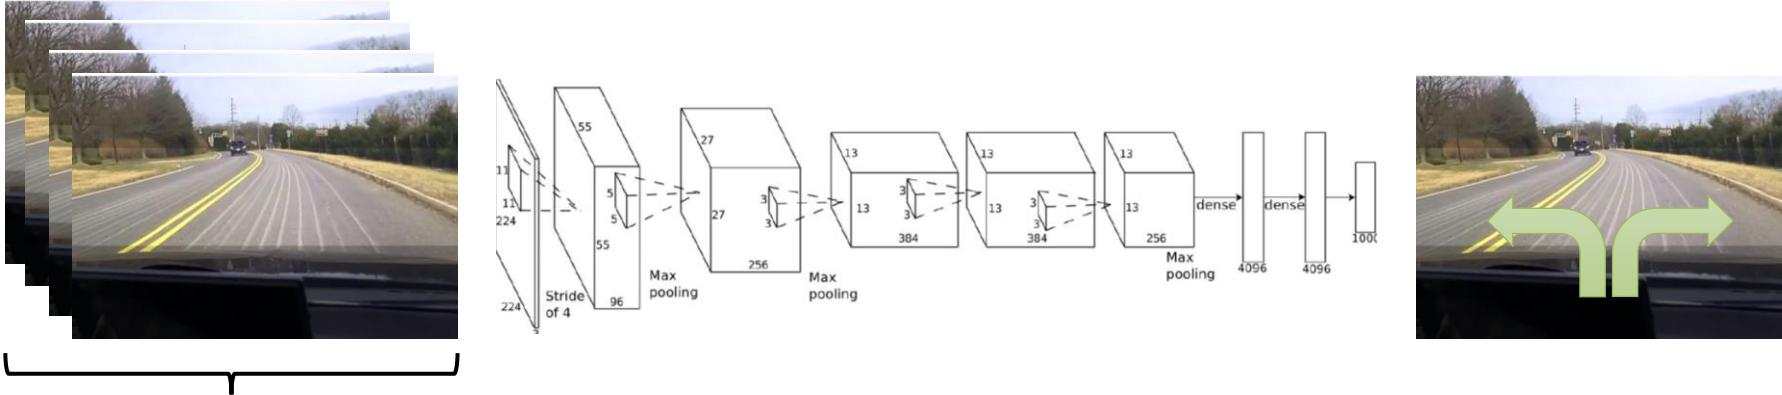
$$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_1, \dots, \mathbf{o}_t)$$

behavior depends on
all past observations

If we see the same thing
twice, we do the same thing
twice, regardless of what
happened before

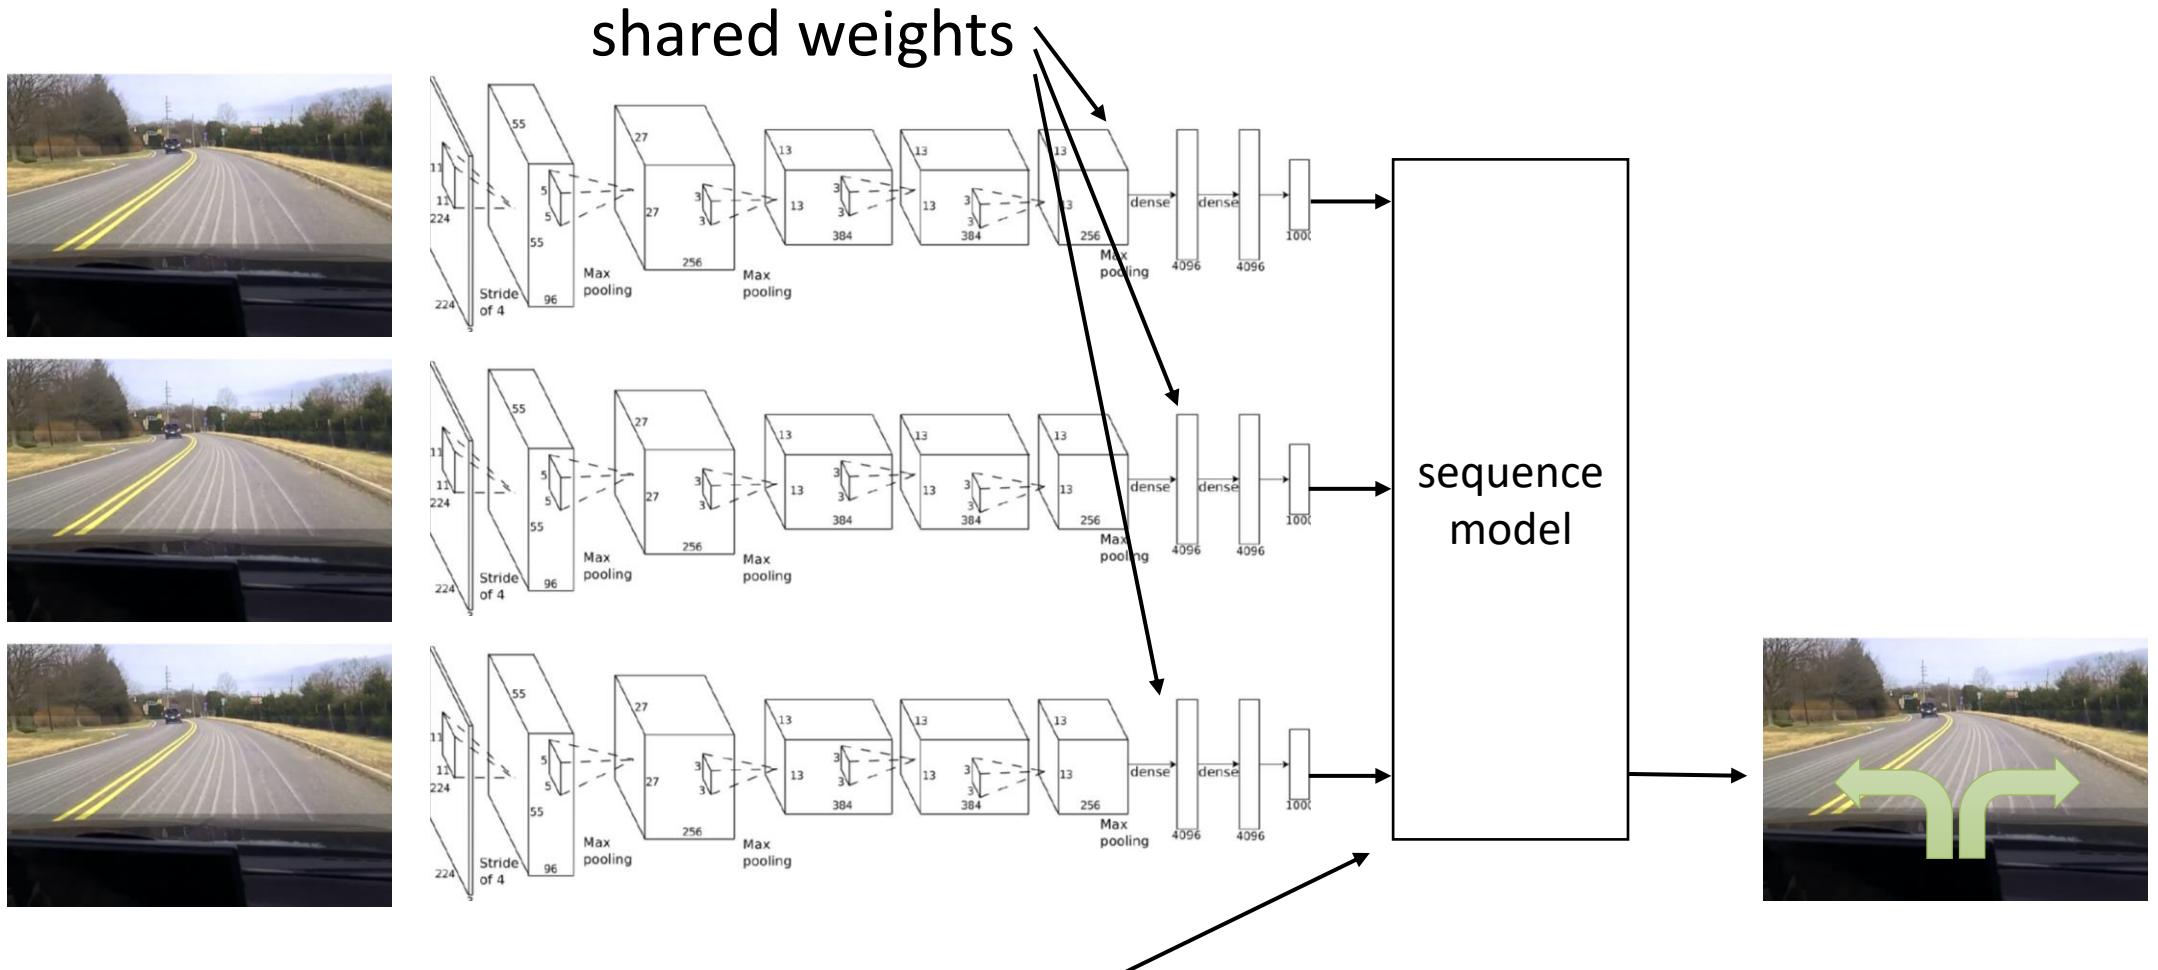
Often very unnatural for
human demonstrators

How can we use the whole history?

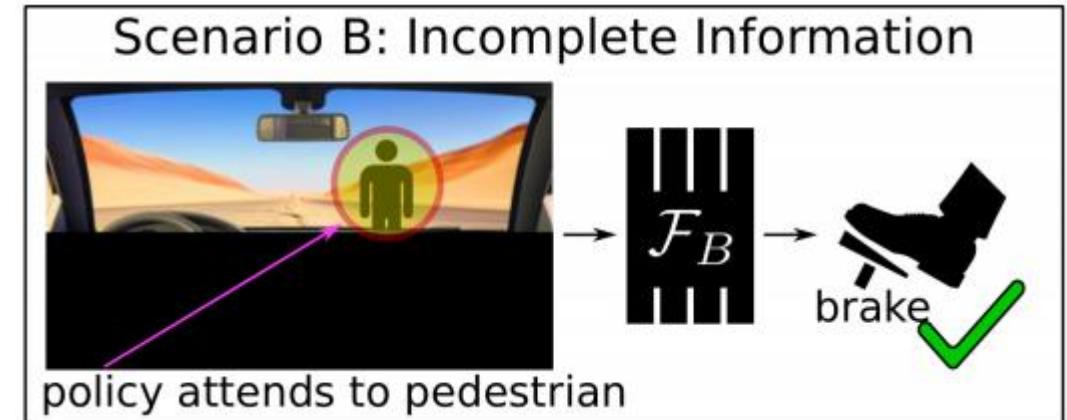
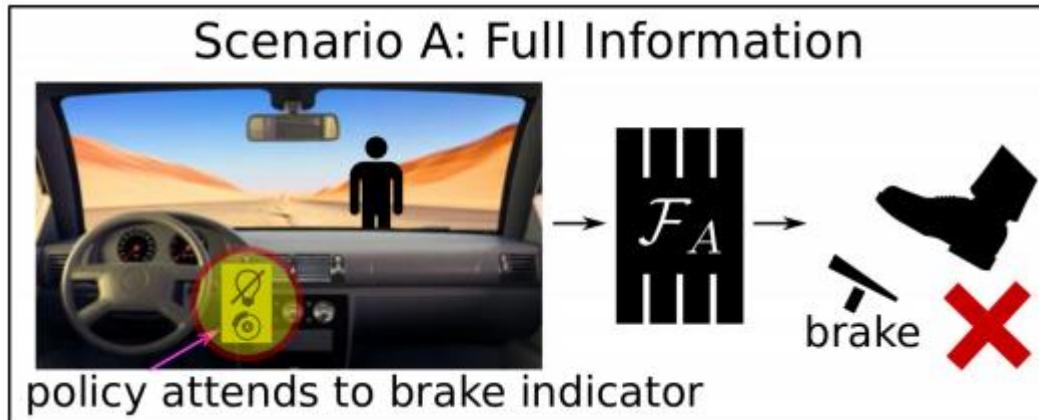


variable number of frames,
too many weights

How can we use the whole history?



Aside: why might this work poorly?



“causal confusion”

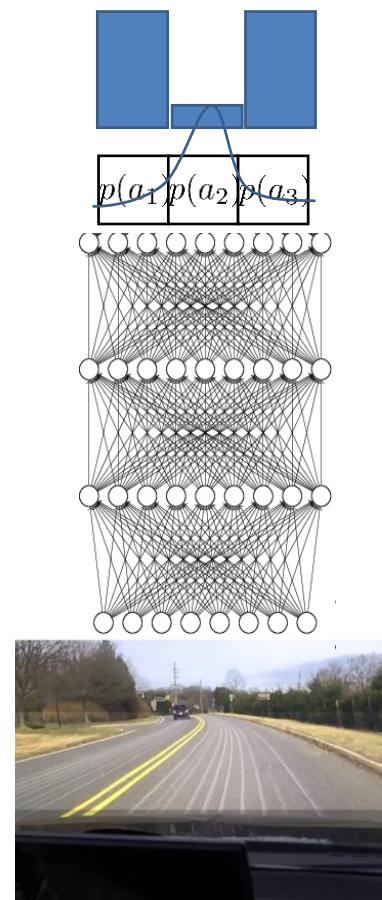
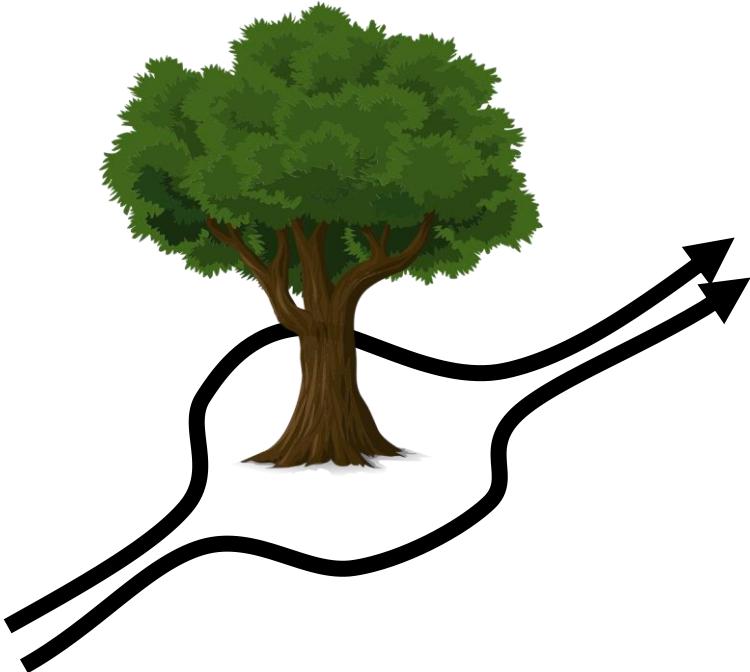
see: de Haan et al., “Causal Confusion in Imitation Learning”

Question 1: Does including history mitigate causal confusion?

Question 2: Can DAgger mitigate causal confusion?

Why might we fail to fit the expert?

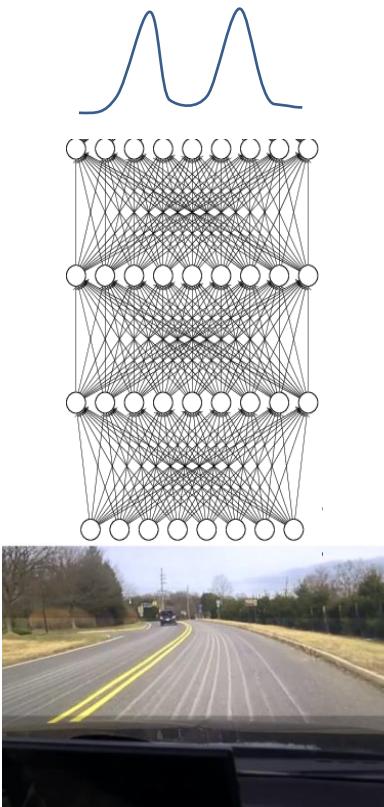
1. Non-Markovian behavior
2. Multimodal behavior



1. More expressive continuous distributions
2. Discretization with high-dimensional action spaces



Expressive continuous distributions



Quite a few options, many ways to make things work:

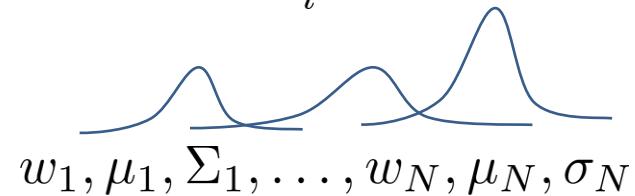
1. mixture of Gaussians
2. latent variable models
3. diffusion models

Expressive continuous distributions

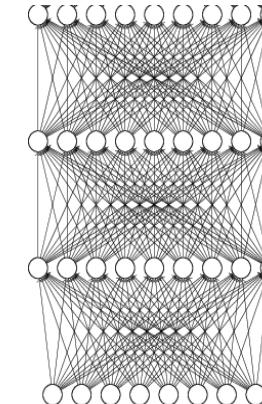


1. mixture of Gaussians
2. latent variable models
3. diffusion models

$$\pi(\mathbf{a}|\mathbf{o}) = \sum_i w_i \mathcal{N}(\mu_i, \Sigma_i)$$



$w_1, \mu_1, \Sigma_1, \dots, w_N, \mu_N, \sigma_N$

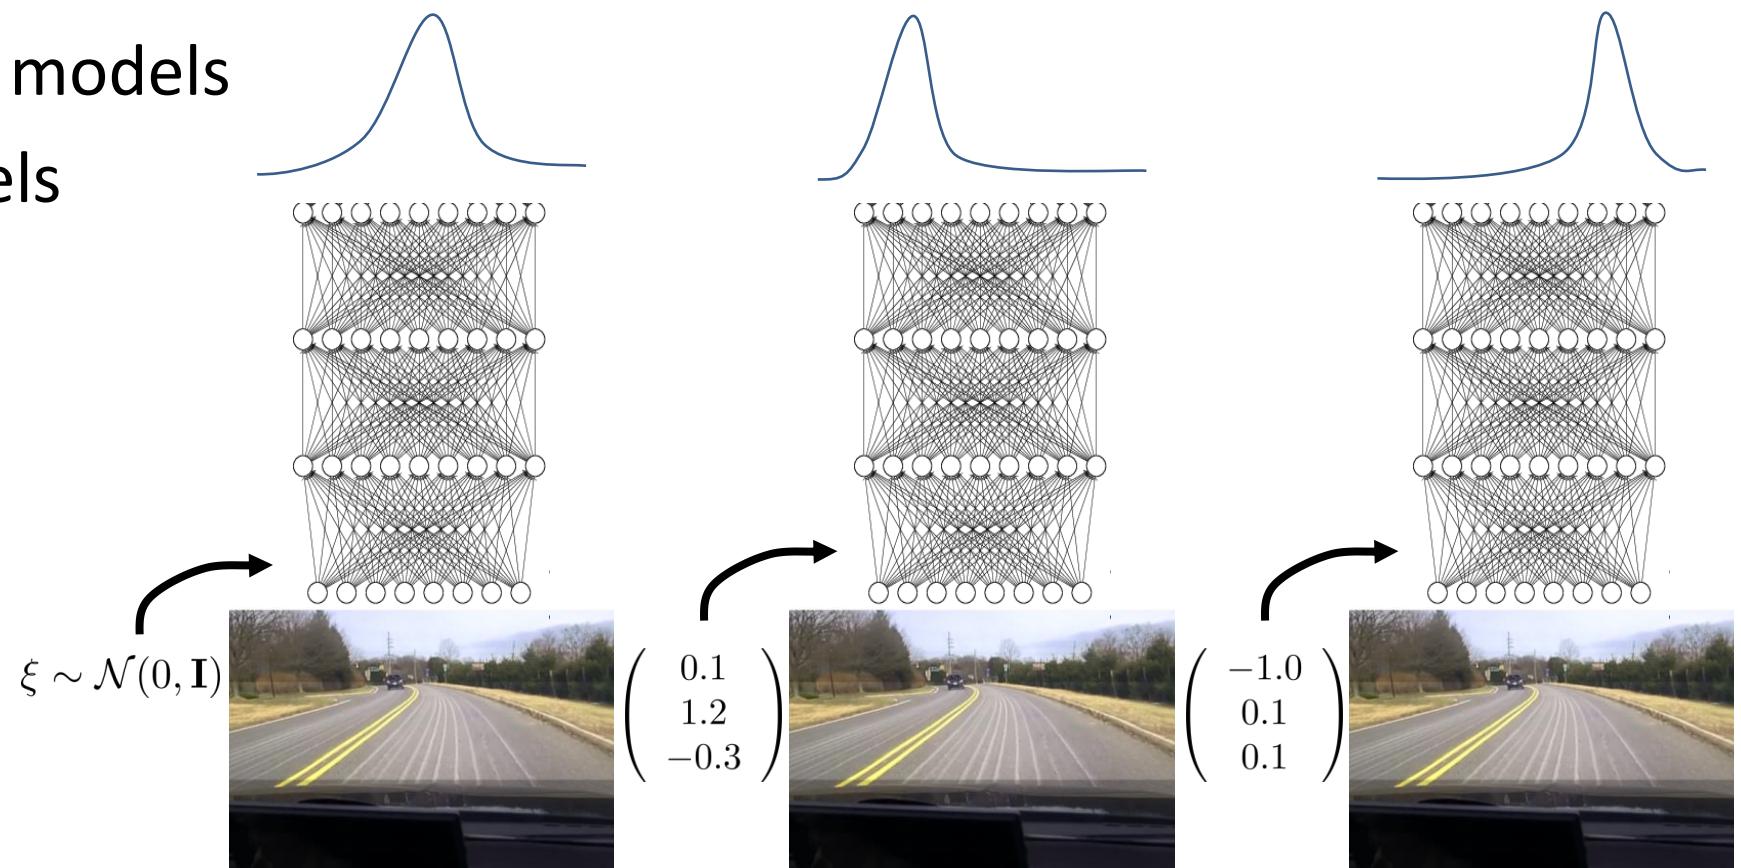


Expressive continuous distributions

- 1. mixture of Gaussians
- 2. latent variable models
- 3. diffusion models

The most widely used type of model of this sort is the (conditional) variational autoencoder

We'll learn about such models later in the course



Expressive continuous distributions

- 1. mixture of Gaussians
- 2. latent variable models
- 3. diffusion models



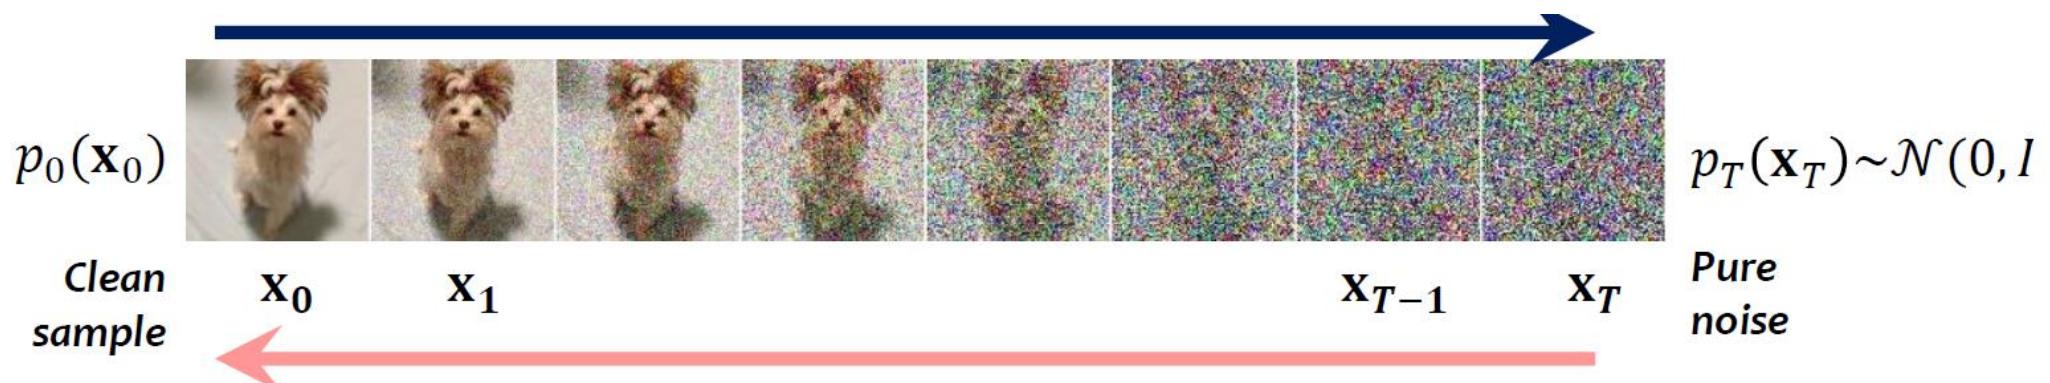
\mathbf{x}_0 = true image

$\mathbf{x}_{i+1} = \mathbf{x}_i + \text{noise}$

Learned network: $f(\mathbf{x}_i) = \mathbf{x}_{i-1}$

(actually use $f(\mathbf{x}_i) = \text{noise}$)

$\mathbf{x}_{i-1} = \mathbf{x}_i - f(\mathbf{x}_i)$



Expressive continuous distributions

1. mixture of Gaussians
2. latent variable models
3. diffusion models



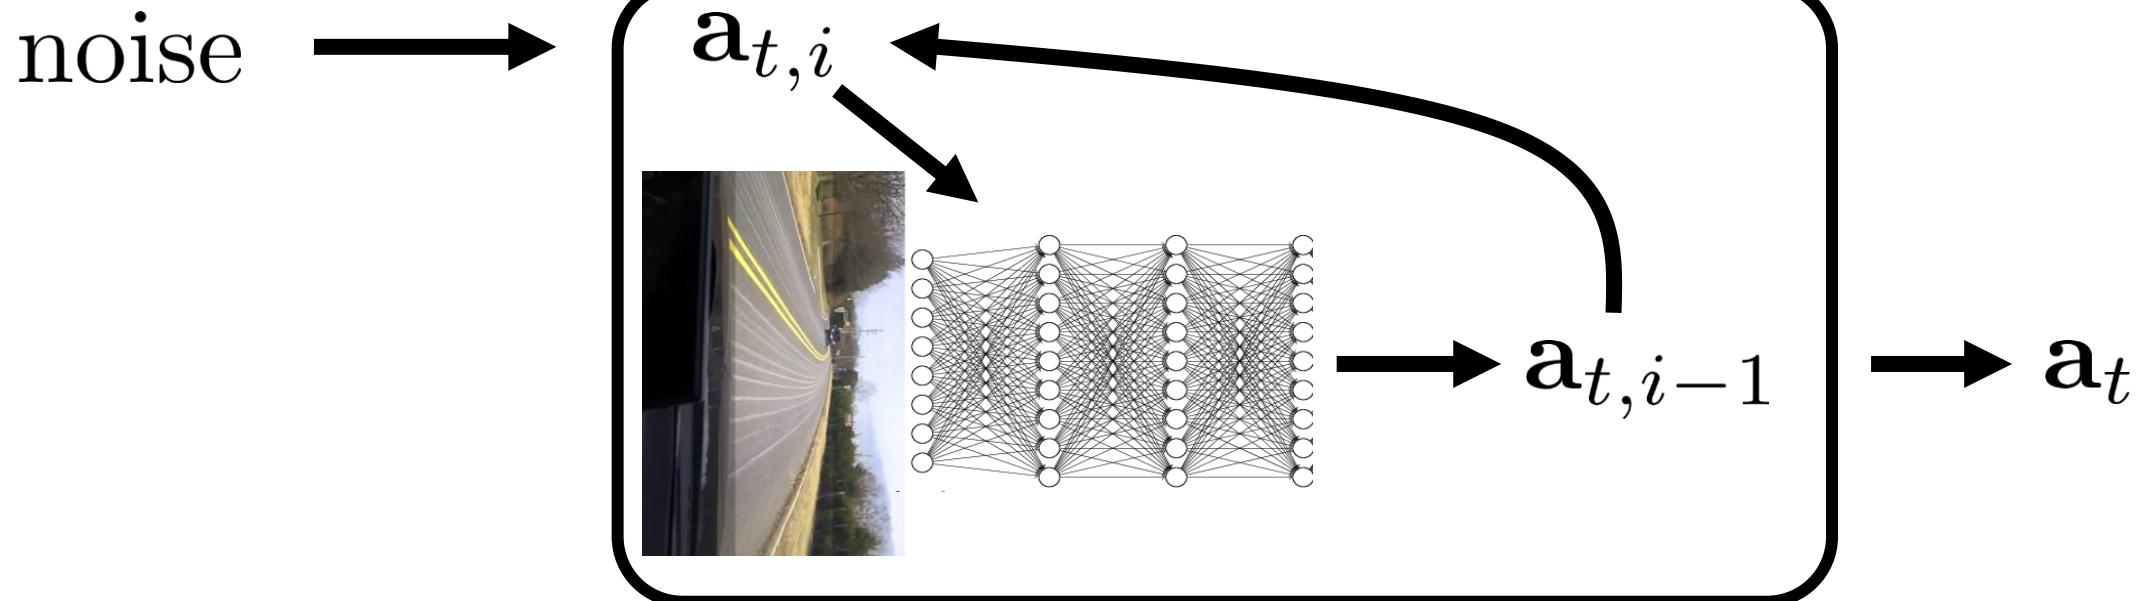
$\mathbf{a}_{t,0}$ = true action

$\mathbf{a}_{t,i+1} = \mathbf{a}_{t,i} + \text{noise}$

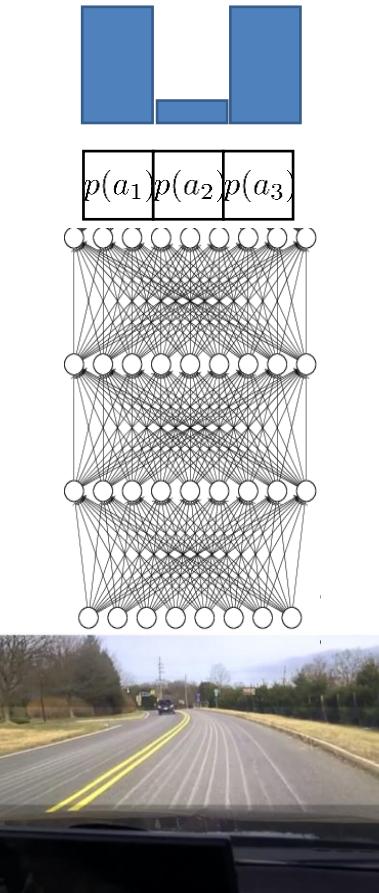
Learned network: $f(\mathbf{s}_t, \mathbf{a}_{t,i}) = \mathbf{a}_{t,i-1}$

(actually use $f(\mathbf{s}_t, \mathbf{a}_{t,i}) = \text{noise}$)

$\mathbf{a}_{t,i-1} = \mathbf{a}_{t,i} - f(\mathbf{s}_t, \mathbf{a}_{t,i})$



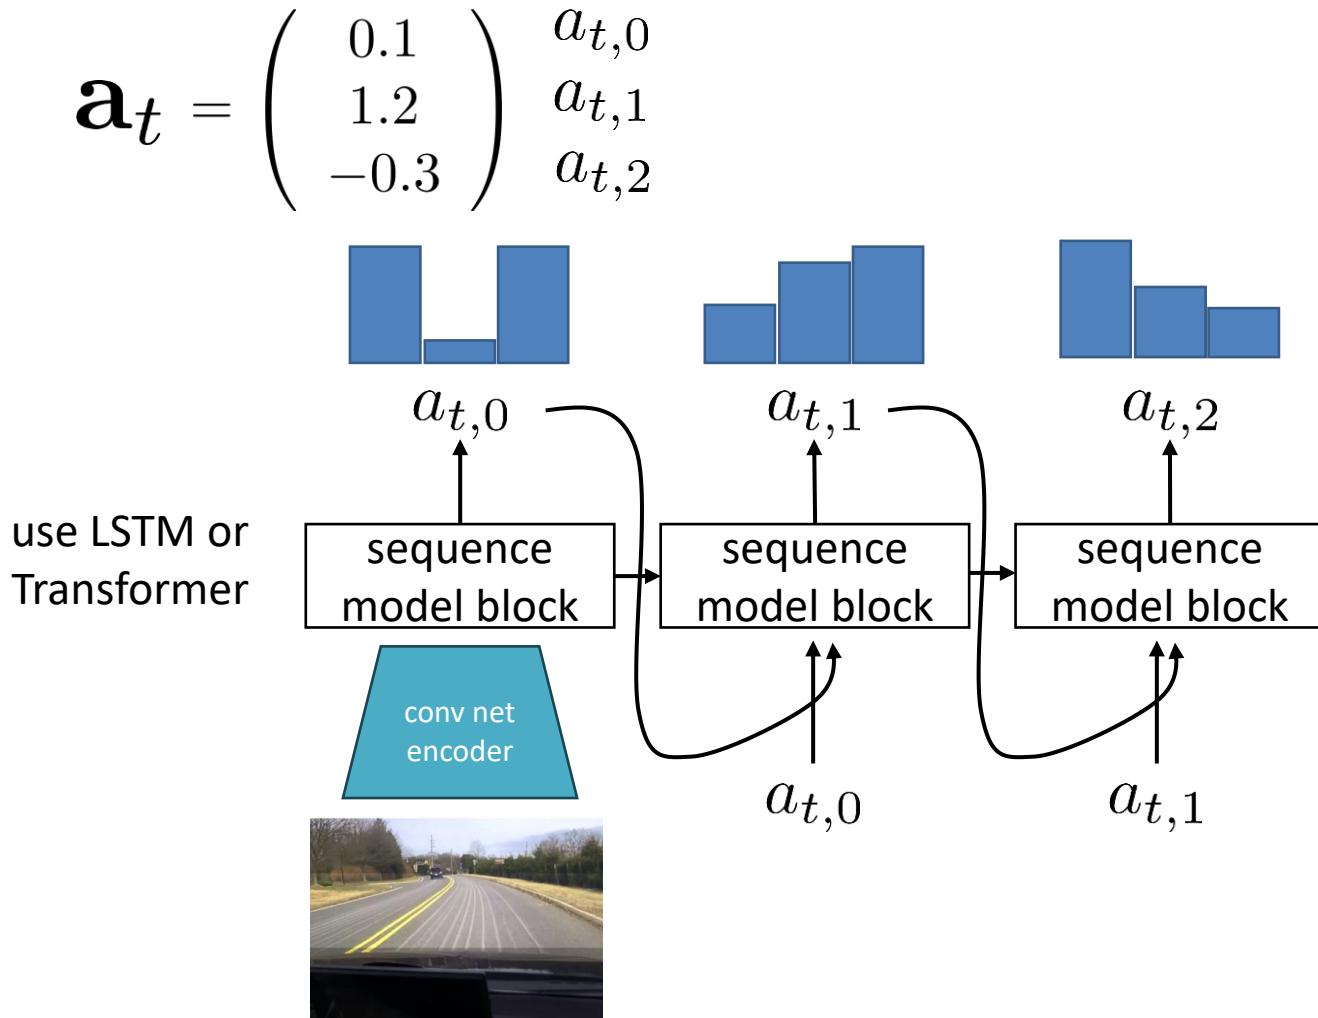
What about discretization?



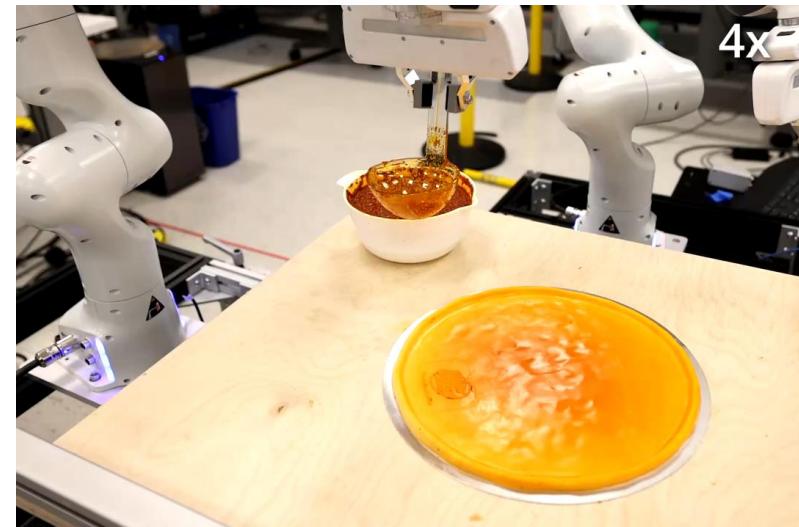
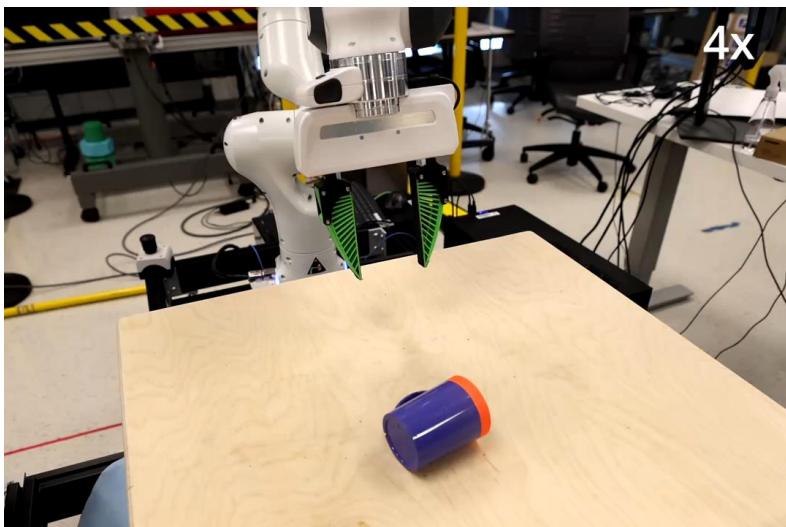
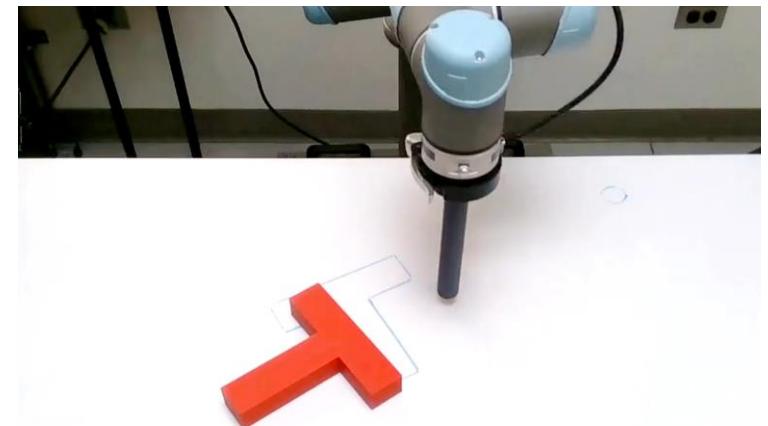
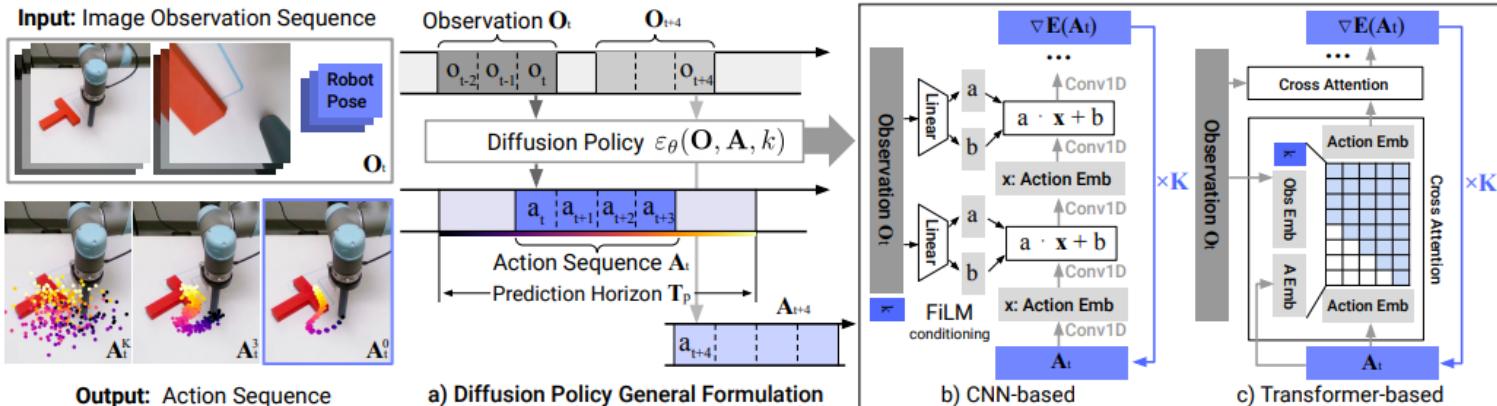
Problem: this is great for 1D actions, but in higher dimensions, discretizing the full space is impractical

Solution: discretize one dimension at a time

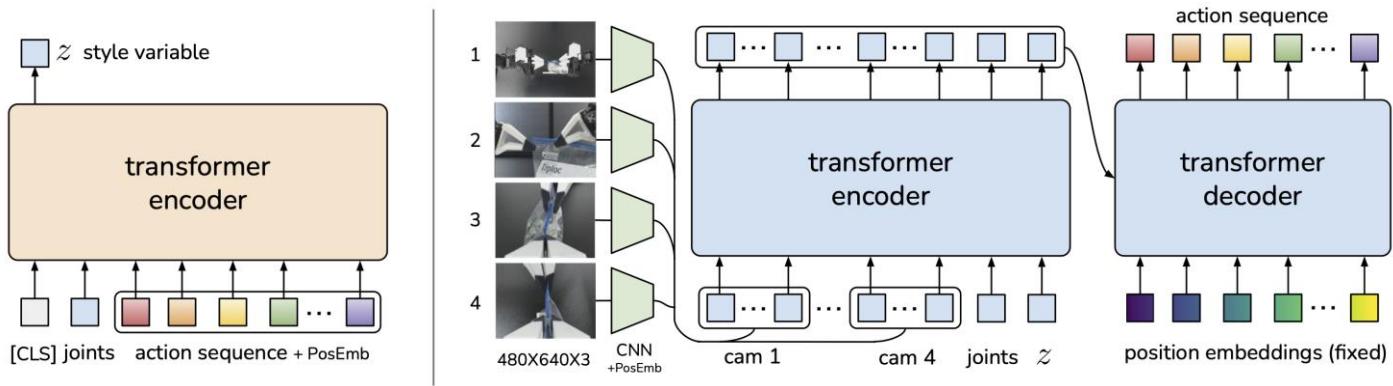
Autoregressive discretization



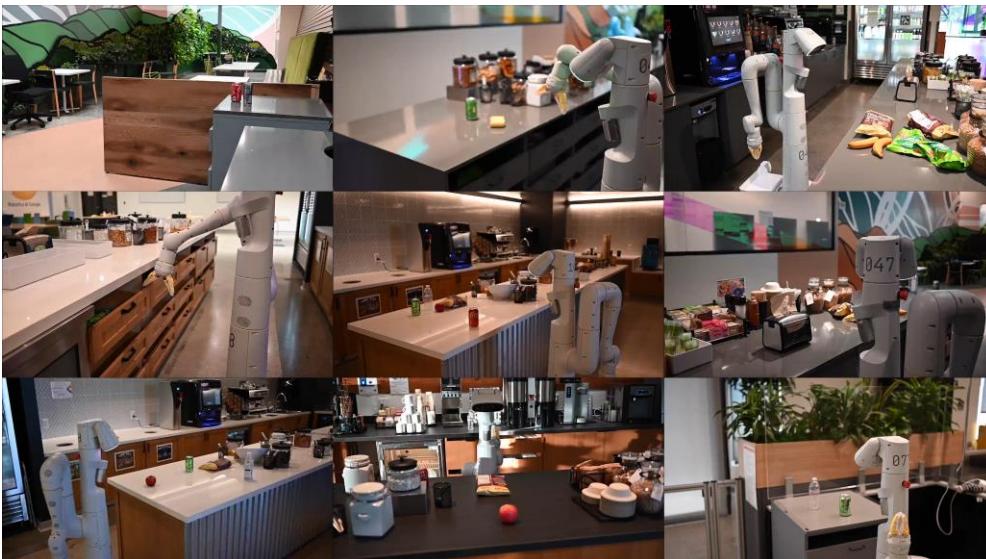
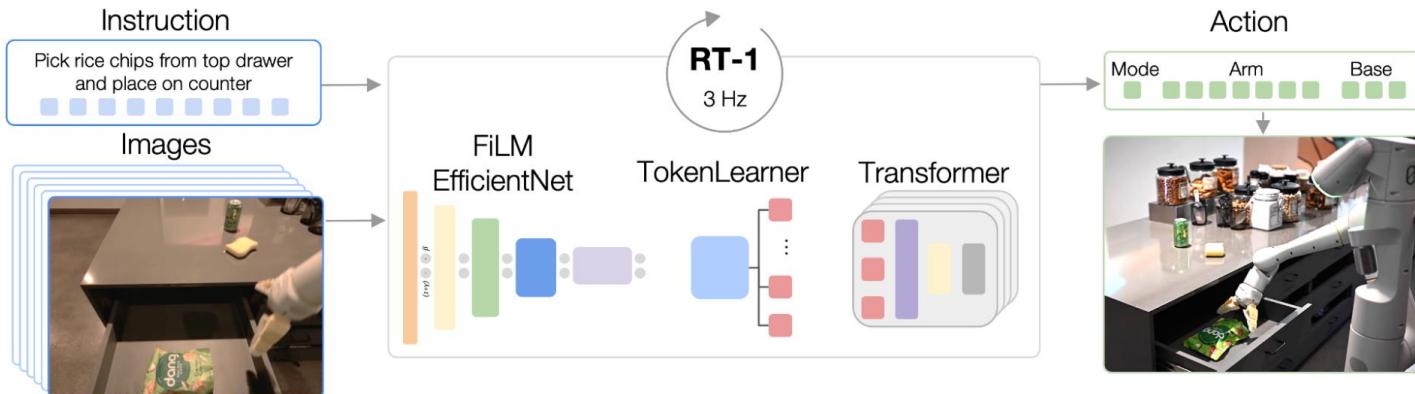
Case study 3: imitation with diffusion models



Case study 4: imitation with latent variables



Case study 5: imitation with Transformers



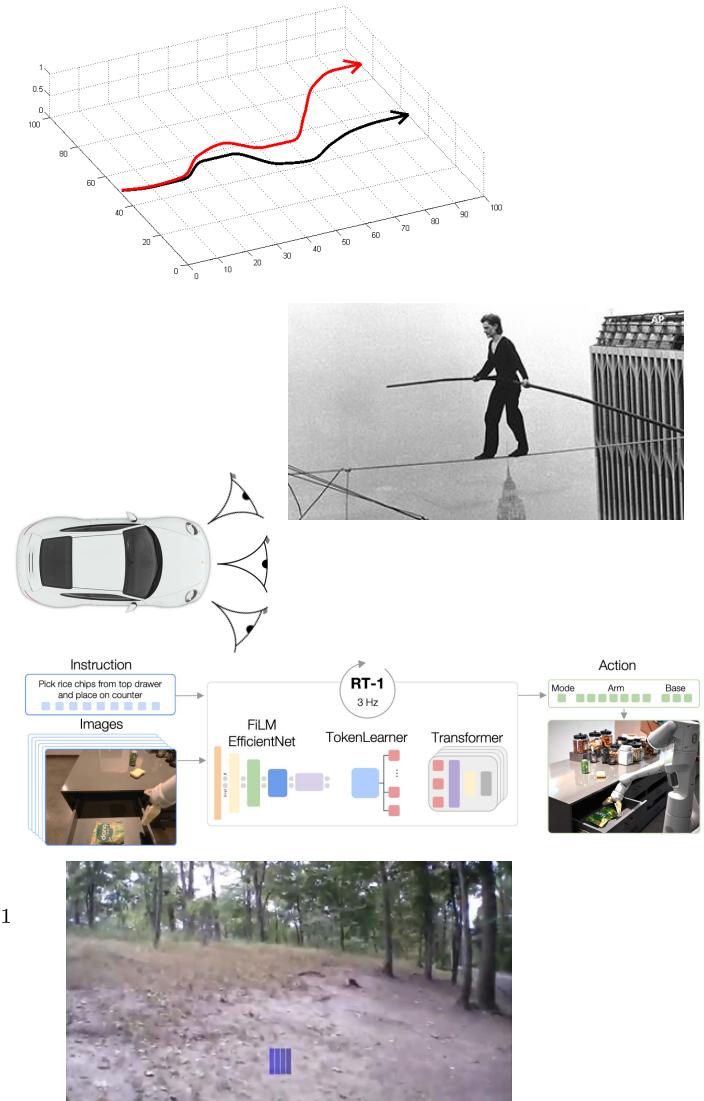
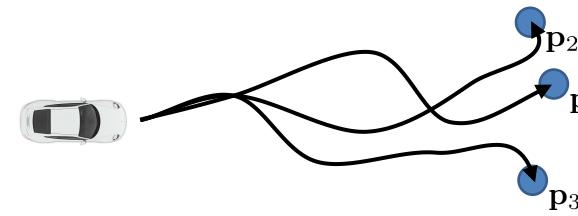
RT-1: Robotics Transformer
for Real-World Control at Scale

 Robotics at Google
 Everyday Robots
Google Research

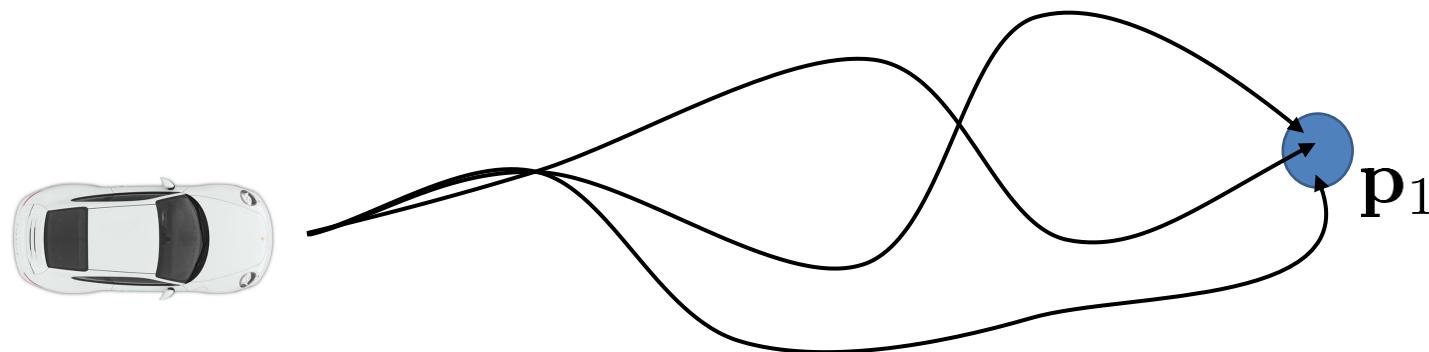
<https://robotics-transformer.github.io/>

Where are we...

- Imitation learning via behavioral cloning is not guaranteed to work
 - This is **different** from supervised learning
 - The reason: i.i.d. assumption does not hold!
- We can formalize **why** this is and do a bit of theory
- **We can address the problem in a few ways:**
 - Be smart about how we collect (and augment) our data
 - Use very powerful models that make very few mistakes
 - **Use multi-task learning**
 - Change the algorithm (Dagger)

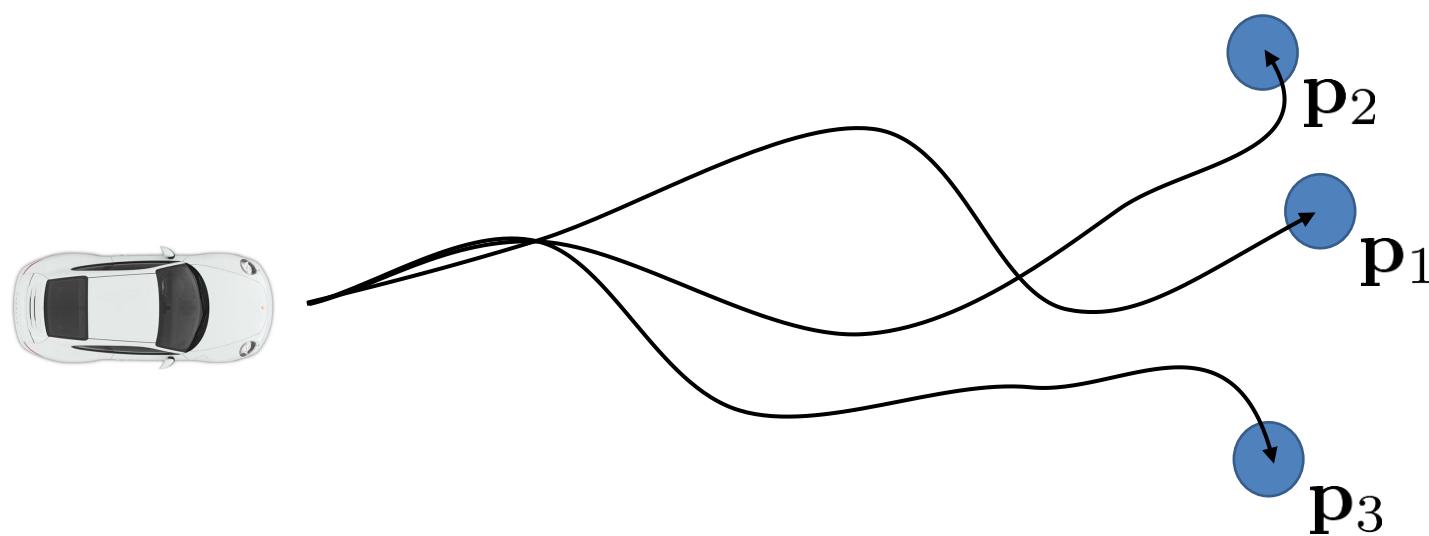


Does learning many tasks become easier?



$$\pi_{\theta}(\mathbf{a}|\mathbf{s})$$

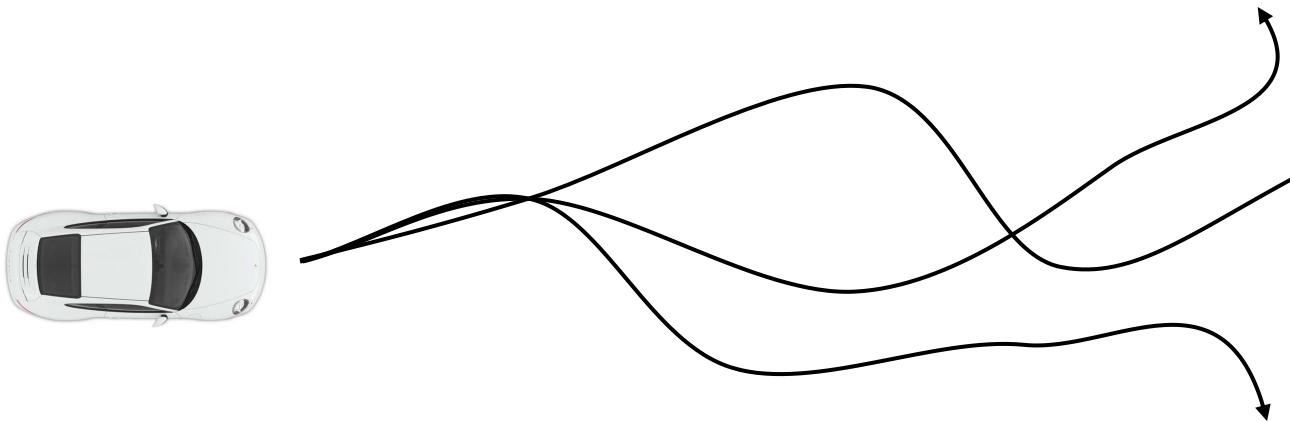
policy for reaching \mathbf{p}_1



$$\pi_{\theta}(\mathbf{a}|\mathbf{s}, \mathbf{p})$$

policy for reaching *any* \mathbf{p}

Goal-conditioned behavioral cloning



training time:

demo 1: $\{s_1, a_t, \dots, s_{T-1}, a_{T-1}, s_T\}$ ← successful demo for reaching s_T

demo 2: $\{s_1, a_t, \dots, s_{T-1}, a_{T-1}, s_T\}$ learn $\pi_\theta(a|s, g)$ ← goal state

demo 3: $\{s_1, a_t, \dots, s_{T-1}, a_{T-1}, s_T\}$

We see distributional shift in **two** places here!

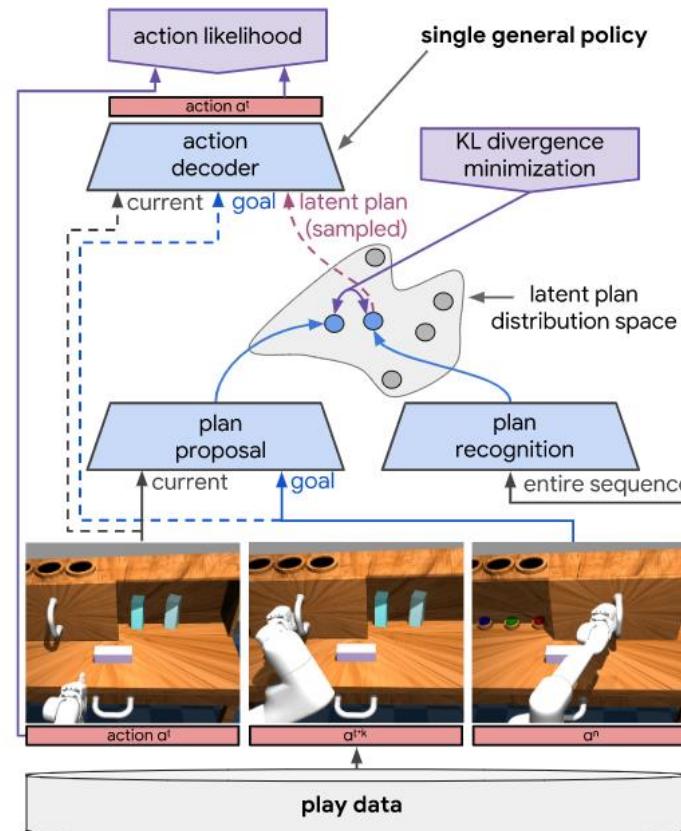
for each demo $\{s_1^i, a_1^i, \dots, s_{T-1}^i, a_{T-1}^i, s_T^i\}$

Can you figure out what the second place is?

maximize $\log \pi_\theta(a_t^i | s_t^i, g = s_T^i)$

Learning Latent Plans from Play

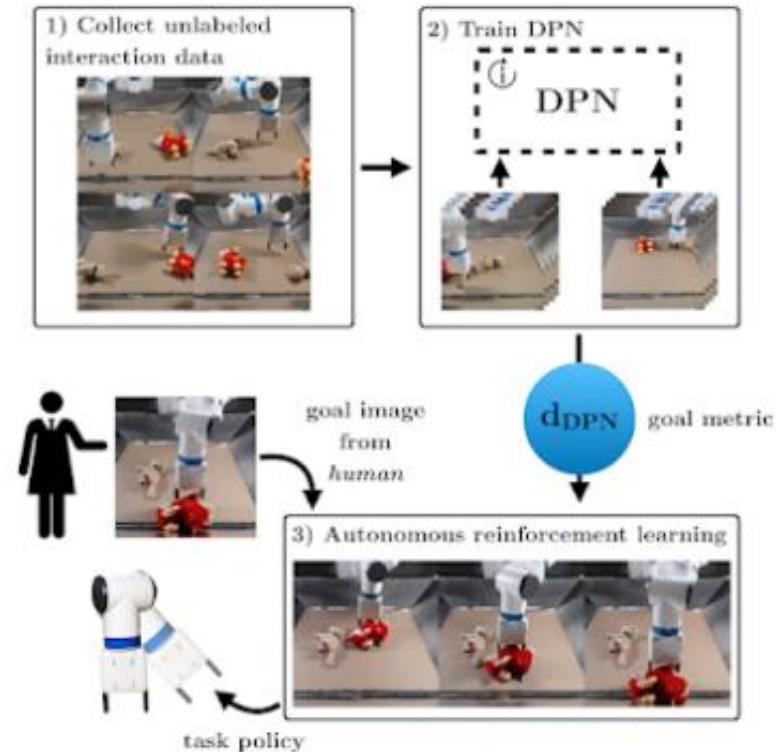
COREY LYNCH MOHI KHANSARI TED XIAO VIKASH KUMAR JONATHAN TOMPSON SERGEY LEVINE PIERRE SERMANET
Google Brain Google X Google Brain Google Brain Google Brain Google Brain Google Brain



Unsupervised Visuomotor Control through Distributional Planning Networks

Tianhe Yu, Gleb Shevchuk, Dorsa Sadigh, Chelsea Finn

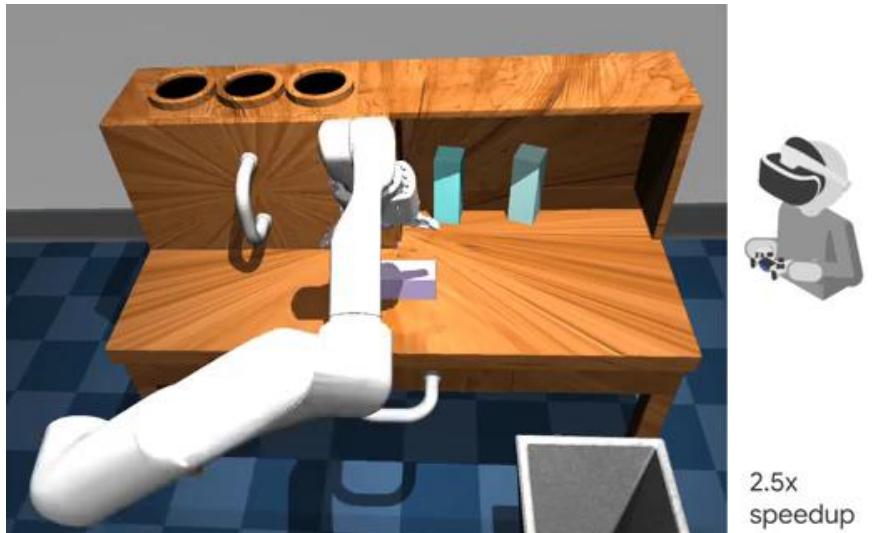
Stanford University



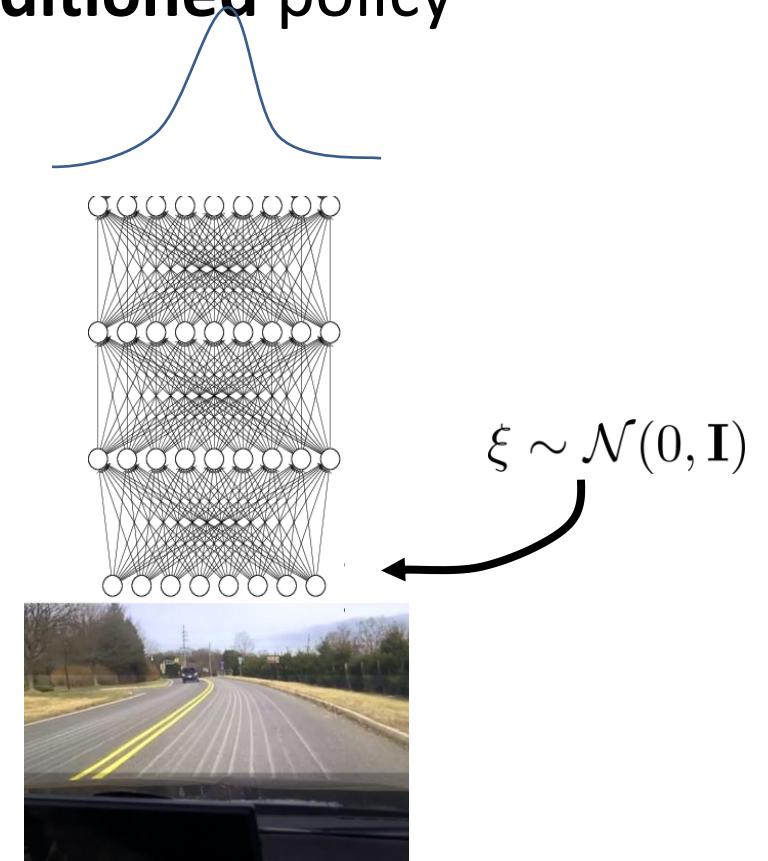
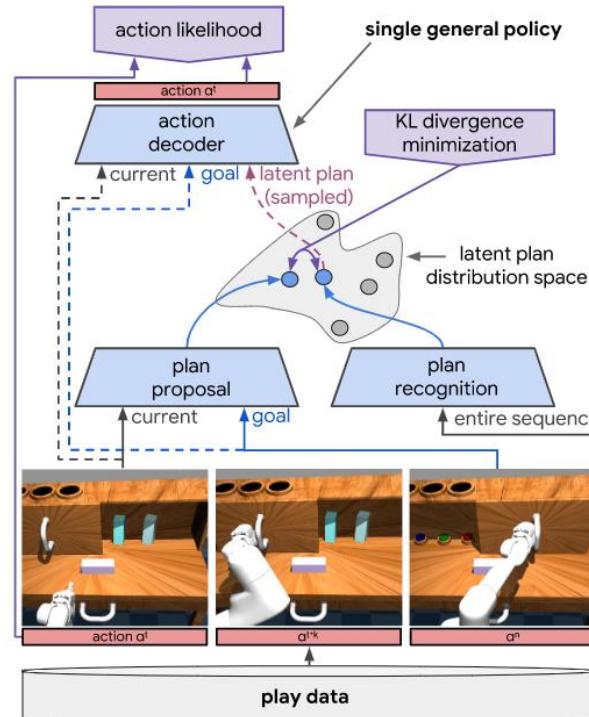
Learning Latent Plans from Play

COREY LYNCH MOHI KHANSARI TED XIAO VIKASH KUMAR JONATHAN TOMPSON SERGEY LEVINE PIERRE SERMANET
Google Brain Google X Google Brain Google Brain Google Brain Google Brain Google Brain

1. Collect data



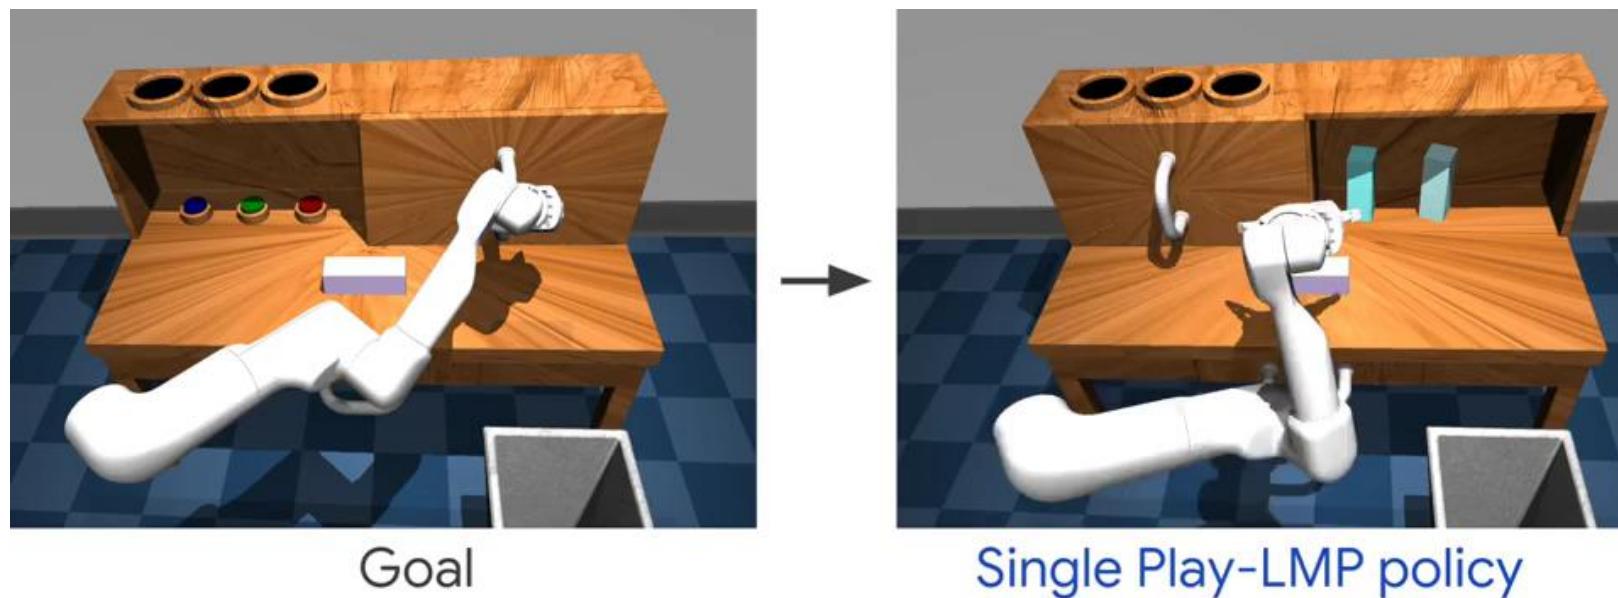
2. Train goal conditioned policy



Learning Latent Plans from Play

COREY LYNCH MOHI KHANSARI TED XIAO VIKASH KUMAR JONATHAN TOMPSON SERGEY LEVINE PIERRE SERMANET
Google Brain Google X Google Brain Google Brain Google Brain Google Brain Google Brain

3. Reach goals



Going beyond just imitation?

Learning to Reach Goals via Iterated Supervised Learning

Dibya Ghosh*
UC Berkeley

Abhishek Gupta*
UC Berkeley

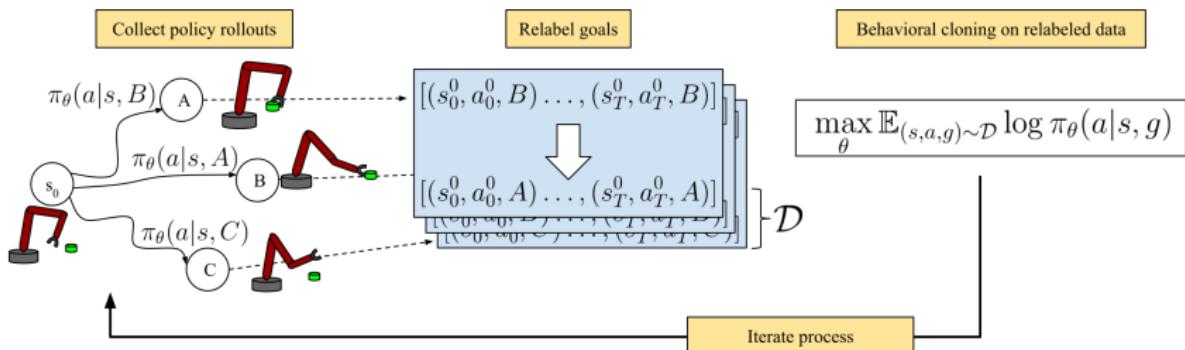
Ashwin Reddy
UC Berkeley

Justin Fu
UC Berkeley

Coline Devin
UC Berkeley

Benjamin Eysenbach
Carnegie Mellon University

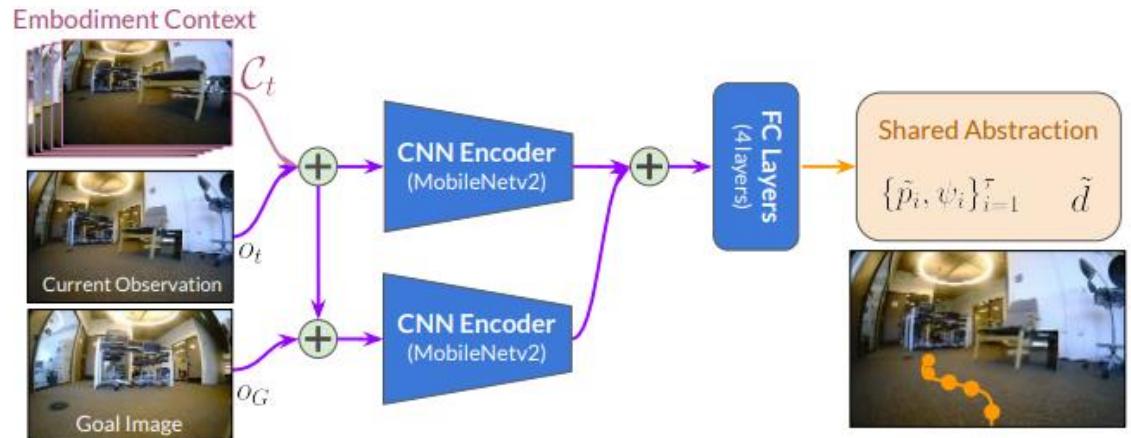
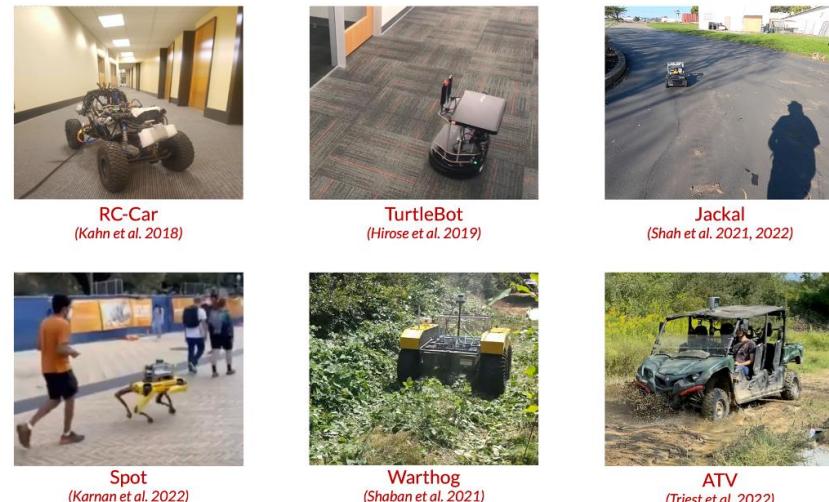
Sergey Levine
UC Berkeley



- Start with a **random** policy
- Collect data with **random** goals
- Treat this data as “demonstrations” for the goals that were reached
- Use this to improve the policy
- Repeat

Goal-conditioned BC at a huge scale

Dataset	Platform	Speed	Amt.	Environment
1 GoStanford [26]	TurtleBot2	0.5m/s	14h	office
2 RECON [32]	Jackal	1m/s	25h	off-road
3 CoryHall [35]	RC Car	1.2m/s	2h	hallways
4 Berkeley [33]	Jackal	2m/s	4h	suburban
5 SCAND-S [36]	Spot	1.5m/s	8h	sidewalks
6 SCAND-J [36]	Jackal	2m/s	1h	sidewalks
7 Seattle [37]	Warthog	5m/s	1h	off-road
8 TartanDrive [38]	ATV	10m/s	5h	off-road
Ours		60h		



Also related (for later...)

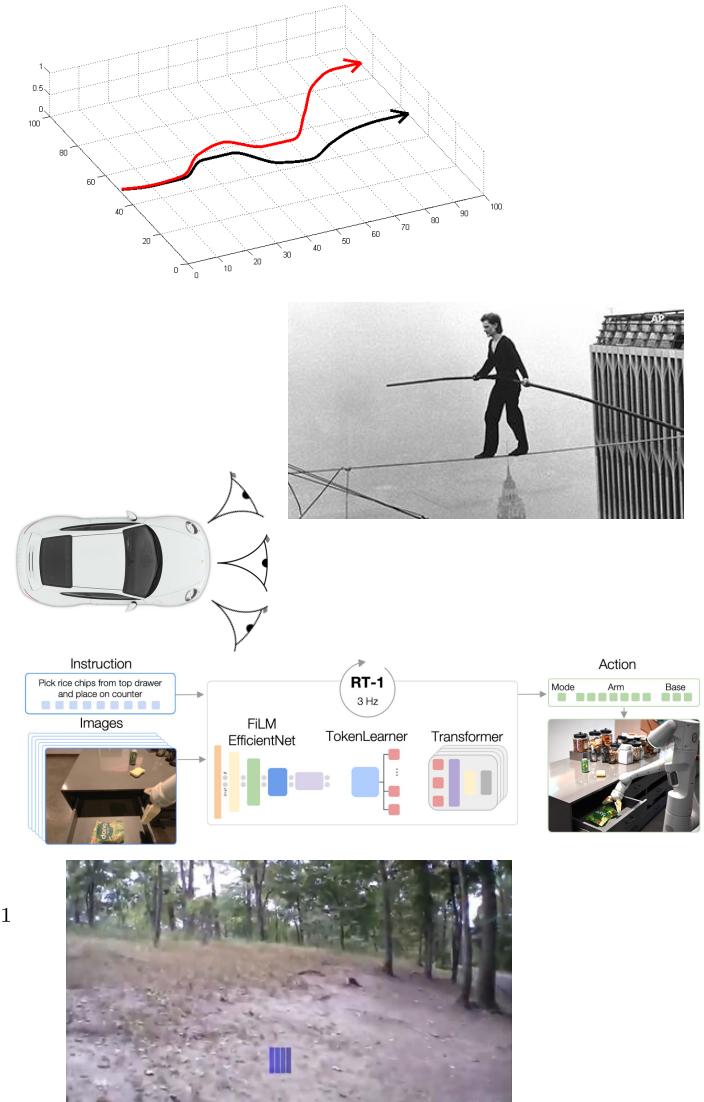
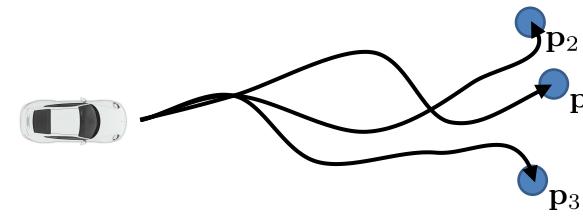
Hindsight Experience Replay

Marcin Andrychowicz*, **Filip Wolski**, **Alex Ray**, **Jonas Schneider**, **Rachel Fong**,
Peter Welinder, **Bob McGrew**, **Josh Tobin**, **Pieter Abbeel†**, **Wojciech Zaremba†**
OpenAI

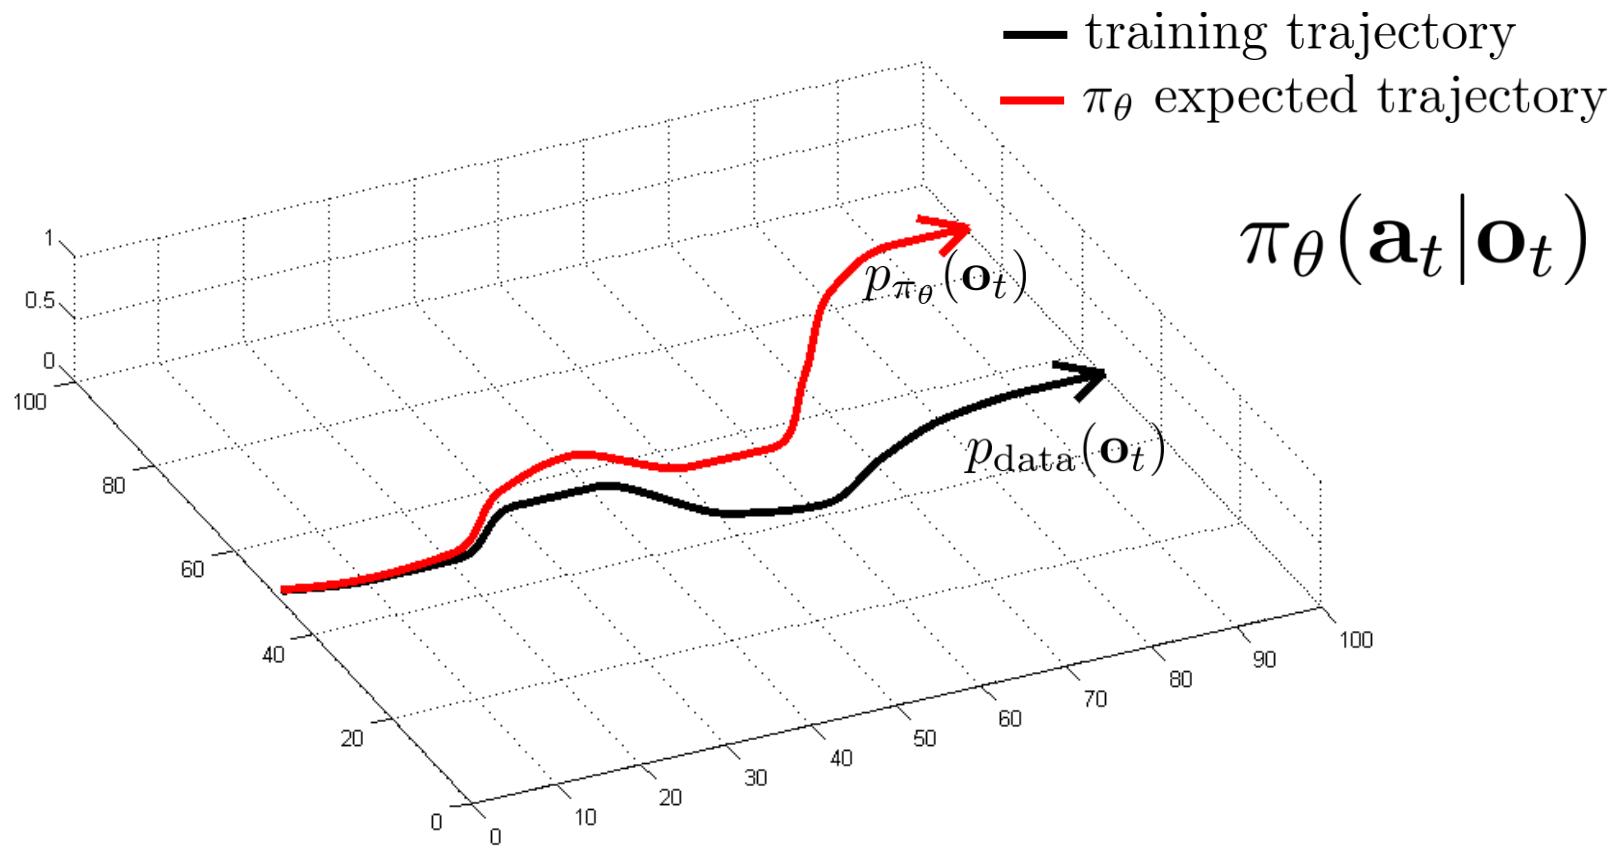
- Similar principle but with reinforcement learning
- This will make more sense later once we cover off-policy value-based RL algorithms
- Worth mentioning because this idea has been used widely outside of imitation (and was arguably first proposed there)

Where are we...

- Imitation learning via behavioral cloning is not guaranteed to work
 - This is **different** from supervised learning
 - The reason: i.i.d. assumption does not hold!
- We can formalize **why** this is and do a bit of theory
- **We can address the problem in a few ways:**
 - Be smart about how we collect (and augment) our data
 - Use very powerful models that make very few mistakes
 - Use multi-task learning
 - Change the algorithm (DAgger)



Can we make it work more often?



can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

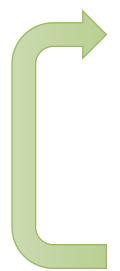
idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{a}_t | \mathbf{o}_t)$

but need labels \mathbf{a}_t !

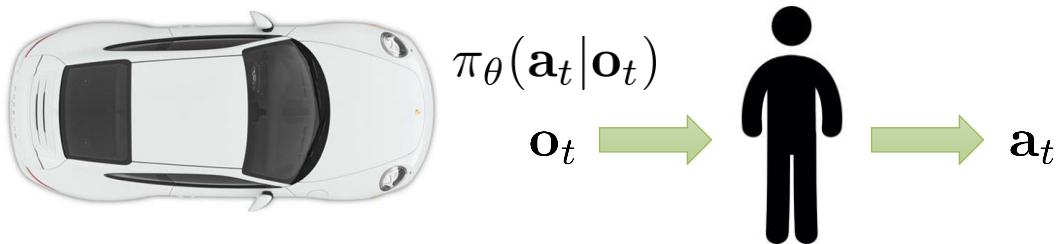
- 
1. train $\pi_\theta(\mathbf{a}_t | \mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{a}_1, \dots, \mathbf{o}_N, \mathbf{a}_N\}$
 2. run $\pi_\theta(\mathbf{a}_t | \mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
 3. Ask human to label \mathcal{D}_π with actions \mathbf{a}_t
 4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

DAgger Example



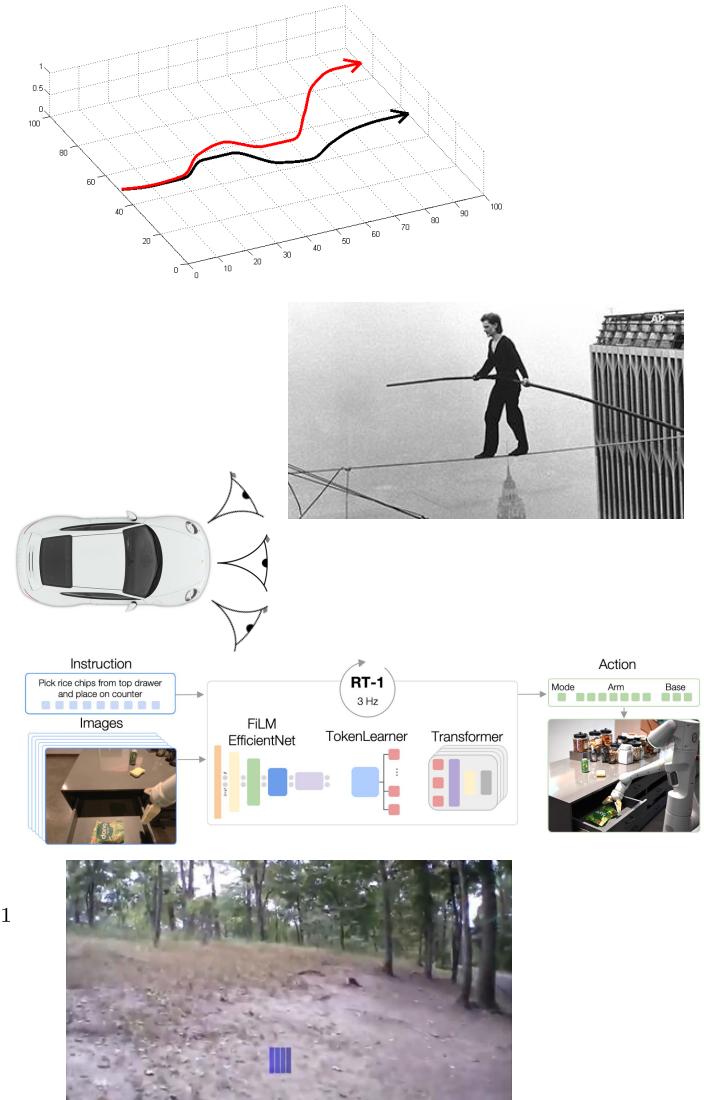
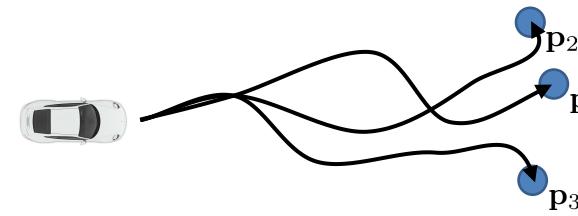
What's the problem?

1. train $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{a}_1, \dots, \mathbf{o}_N, \mathbf{a}_N\}$
2. run $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
3. Ask human to label \mathcal{D}_π with actions \mathbf{a}_t
4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$



Recap

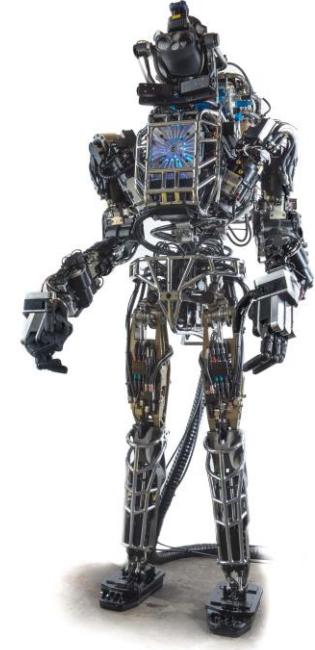
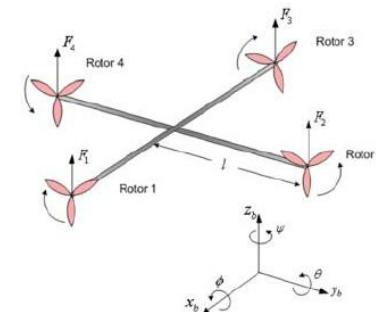
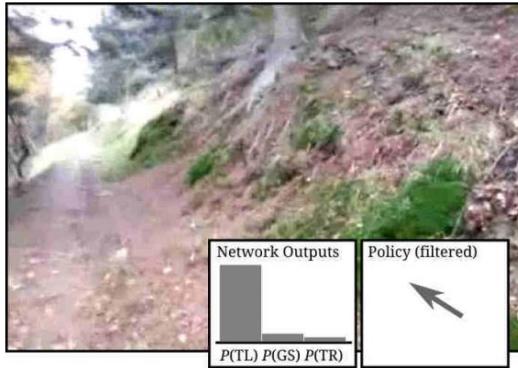
- Imitation learning via behavioral cloning is not guaranteed to work
 - This is **different** from supervised learning
 - The reason: i.i.d. assumption does not hold!
- We can formalize **why** this is and do a bit of theory
- We can address the problem in a few ways:
 - Be smart about how we collect (and augment) our data
 - Use very powerful models that make very few mistakes
 - Use multi-task learning
 - Change the algorithm (DAgger)



Cost functions and reward functions,
a preview of what comes next

Imitation learning: what's the problem?

- Humans need to provide data, which is typically finite
 - Deep learning works best when data is plentiful
- Humans are not good at providing some kinds of actions



- Humans can learn autonomously; can our machines do the same?
 - Unlimited data from own experience
 - Continuous self-improvement

Terminology & notation



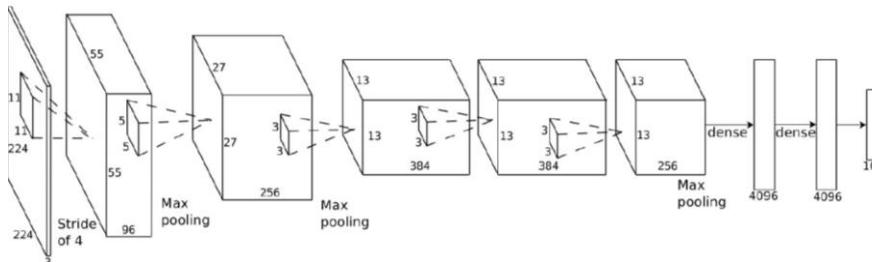
\mathbf{o}_t



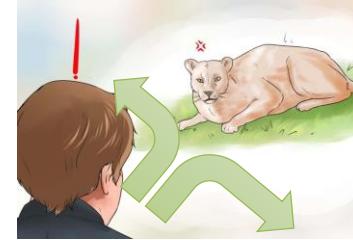
\mathbf{s}_t – state

\mathbf{o}_t – observation

\mathbf{a}_t – action



$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)$



\mathbf{a}_t



$c(\mathbf{s}_t, \mathbf{a}_t)$ – cost function

$r(\mathbf{s}_t, \mathbf{a}_t)$ – reward function

$$\min_{\theta} E_{\mathbf{a}_t \sim \pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t), \mathbf{s}_t \sim p(\mathbf{s}_t | \mathbf{s}_{t-1}, \mathbf{a}_{t-1})} \left[\sum_t \delta(s_t, a_t) r(s_t, a_t) + \beta \mathbb{E}_{\mathbf{s}'_t \sim p(\mathbf{s}'_t | \mathbf{s}_t, \mathbf{a}_t)} [c(\mathbf{s}'_t, \mathbf{a}_t)] \right]$$

[
 $\sum_t \delta(s_t, a_t) r(s_t, a_t)$
 + $\beta \mathbb{E}_{\mathbf{s}'_t \sim p(\mathbf{s}'_t | \mathbf{s}_t, \mathbf{a}_t)} [c(\mathbf{s}'_t, \mathbf{a}_t)]$]
]

by tiger) by tiger)

Aside: notation

s_t – state

a_t – action

$r(s, a)$ – reward function



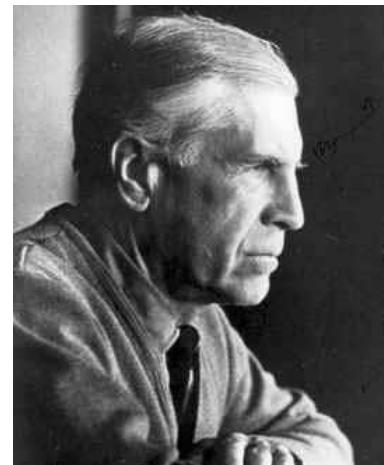
Richard Bellman

x_t – state

u_t – action

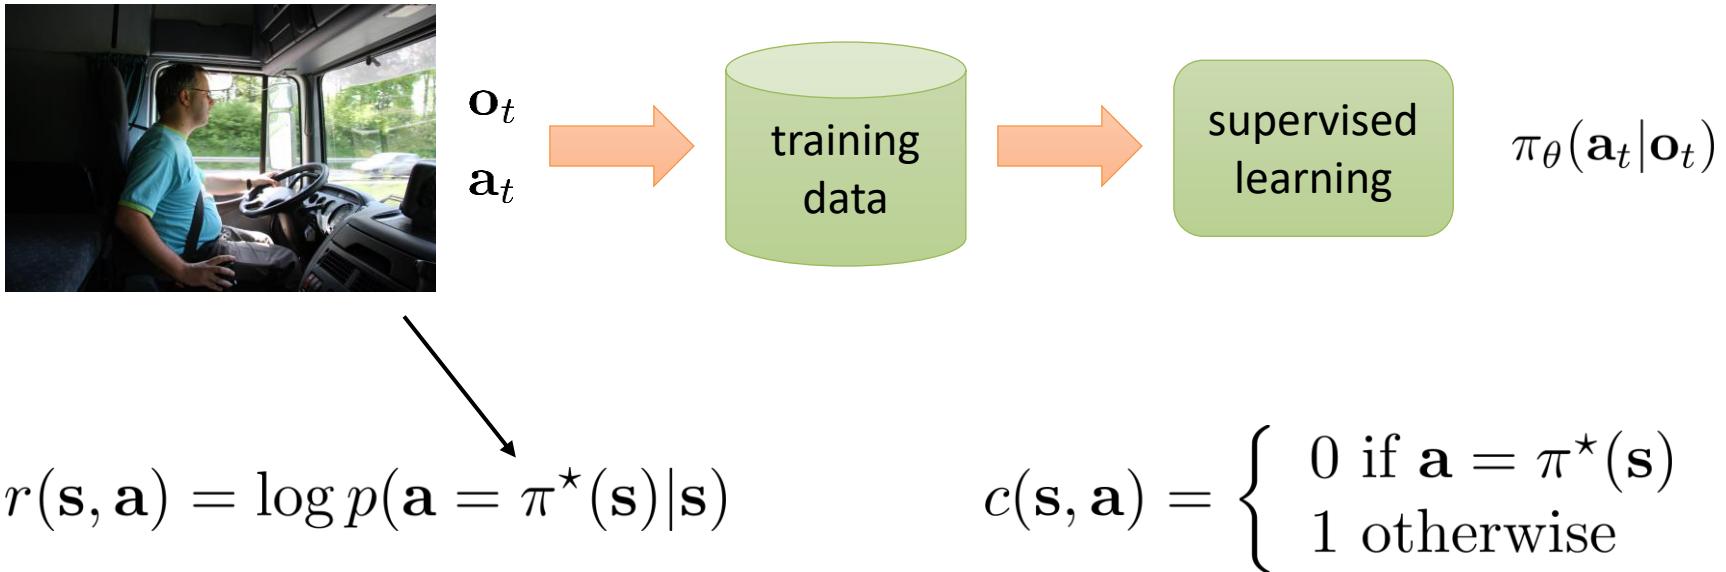
$c(x, u)$ – cost function

$$r(s, a) = -c(x, u)$$



Lev Pontryagin

A cost function for imitation?



Goal: minimize $E_{\mathbf{s}_t \sim p_{\pi_\theta}(\mathbf{s}_t)}[c(\mathbf{s}_t, \mathbf{a}_t)]$

Goal: maximize $E_{\mathbf{s}_t \sim p_{\pi_\theta}(\mathbf{s}_t)}[r(\mathbf{s}_t, \mathbf{a}_t)]$

Imitation learning algorithms **do** maximize reward when they work well!

For a very particular choice of reward